

# Advanced Regression

## Surprise Housing Case Study

A US-based housing company named Surprise Housing has decided to enter the Australian market. The company uses data analytics to purchase houses at a price below their actual values and flip them on at a higher price. For the same purpose, the company has collected a data set from the sale of houses in Australia.

### **\*\*Problem Statement\*\***

The company is looking at prospective properties to buy to enter the market. It is required to build a regression model using regularisation in order to predict the actual value of the prospective properties and decide whether to invest in them or not.

The company wants to know:

- Which variables are significant in predicting the price of a house, and
- How well those variables describe the price of a house.

Also, determine the optimal value of lambda for ridge and lasso regression.

### **\*\*Business Goal\*\***

To model the price of houses with the available independent variables. This model will then be used by the management to understand how exactly the prices vary with the variables. They can accordingly manipulate the strategy of the firm and concentrate on areas that will yield high returns. Further, the model will be a good way for management to understand the pricing dynamics of a new market.

The solution is divided into the following sections:

- Data understanding and exploration
- Data cleaning
- Data preparation
- Model building and evaluation

## 1. Data Understanding and Exploration

Let's first have a look at the dataset and understand the size, attribute names etc.

```
In [1]: import numpy as np
import pandas as pd
```

```

import matplotlib.pyplot as plt
import seaborn as sns
import datetime
from sklearn import linear_model, metrics
from sklearn.linear_model import LinearRegression
from sklearn.linear_model import Ridge
from sklearn.linear_model import Lasso
from sklearn.model_selection import GridSearchCV
from sklearn.metrics import mean_squared_error, r2_score

import os

# hide warnings
import warnings
warnings.filterwarnings('ignore')

pd.set_option('display.max_rows', 500)

```

```

In [2]: # reading the dataset
df = pd.read_csv("train.csv")
df.head()

```

```

Out[2]:
```

	Id	MSSubClass	MSZoning	LotFrontage	LotArea	Street	Alley	LotShape	LandContour	Utilities
0	1	60	RL	65.0	8450	Pave	NaN	Reg	Lvl	AllPub
1	2	20	RL	80.0	9600	Pave	NaN	Reg	Lvl	AllPub
2	3	60	RL	68.0	11250	Pave	NaN	IR1	Lvl	AllPub
3	4	70	RL	60.0	9550	Pave	NaN	IR1	Lvl	AllPub
4	5	60	RL	84.0	14260	Pave	NaN	IR1	Lvl	AllPub

5 rows × 81 columns

```

In [3]: # check the shape

df.shape

```

```

Out[3]: (1460, 81)

```

The dataset has 1460 rows and 81 columns.

```

In [4]: df.info()

```

<class 'pandas.core.frame.DataFrame'>

RangeIndex: 1460 entries, 0 to 1459

Data columns (total 81 columns):

#	Column	Non-Null Count	Dtype
0	Id	1460 non-null	int64
1	MSSubClass	1460 non-null	int64
2	MSZoning	1460 non-null	object
3	LotFrontage	1201 non-null	float64
4	LotArea	1460 non-null	int64
5	Street	1460 non-null	object
6	Alley	91 non-null	object
7	LotShape	1460 non-null	object
8	LandContour	1460 non-null	object
9	Utilities	1460 non-null	object
10	LotConfig	1460 non-null	object
11	LandSlope	1460 non-null	object
12	Neighborhood	1460 non-null	object
13	Condition1	1460 non-null	object
14	Condition2	1460 non-null	object
15	BldgType	1460 non-null	object
16	HouseStyle	1460 non-null	object
17	OverallQual	1460 non-null	int64
18	OverallCond	1460 non-null	int64
19	YearBuilt	1460 non-null	int64
20	YearRemodAdd	1460 non-null	int64
21	RoofStyle	1460 non-null	object
22	RoofMatl	1460 non-null	object
23	Exterior1st	1460 non-null	object
24	Exterior2nd	1460 non-null	object
25	MasVnrType	1452 non-null	object
26	MasVnrArea	1452 non-null	float64
27	ExterQual	1460 non-null	object
28	ExterCond	1460 non-null	object
29	Foundation	1460 non-null	object
30	BsmtQual	1423 non-null	object
31	BsmtCond	1423 non-null	object
32	BsmtExposure	1422 non-null	object
33	BsmtFinType1	1423 non-null	object
34	BsmtFinSF1	1460 non-null	int64
35	BsmtFinType2	1422 non-null	object
36	BsmtFinSF2	1460 non-null	int64
37	BsmtUnfSF	1460 non-null	int64
38	TotalBsmtSF	1460 non-null	int64
39	Heating	1460 non-null	object
40	HeatingQC	1460 non-null	object
41	CentralAir	1460 non-null	object
42	Electrical	1459 non-null	object
43	1stFlrSF	1460 non-null	int64
44	2ndFlrSF	1460 non-null	int64
45	LowQualFinSF	1460 non-null	int64
46	GrLivArea	1460 non-null	int64
47	BsmtFullBath	1460 non-null	int64
48	BsmtHalfBath	1460 non-null	int64
49	FullBath	1460 non-null	int64
50	HalfBath	1460 non-null	int64
51	BedroomAbvGr	1460 non-null	int64
52	KitchenAbvGr	1460 non-null	int64
53	KitchenQual	1460 non-null	object
54	TotRmsAbvGrd	1460 non-null	int64

```

55 Functional      1460 non-null  object
56 Fireplaces      1460 non-null  int64
57 FireplaceQu     770 non-null   object
58 GarageType      1379 non-null  object
59 GarageYrBlt     1379 non-null  float64
60 GarageFinish    1379 non-null  object
61 GarageCars      1460 non-null  int64
62 GarageArea      1460 non-null  int64
63 GarageQual      1379 non-null  object
64 GarageCond      1379 non-null  object
65 PavedDrive      1460 non-null  object
66 WoodDeckSF      1460 non-null  int64
67 OpenPorchSF     1460 non-null  int64
68 EnclosedPorch   1460 non-null  int64
69 3SsnPorch       1460 non-null  int64
70 ScreenPorch     1460 non-null  int64
71 PoolArea        1460 non-null  int64
72 PoolQC          7 non-null     object
73 Fence           281 non-null   object
74 MiscFeature      54 non-null    object
75 MiscVal          1460 non-null  int64
76 MoSold          1460 non-null  int64
77 YrSold           1460 non-null  int64
78 SaleType         1460 non-null  object
79 SaleCondition    1460 non-null  object
80 SalePrice        1460 non-null  int64
dtypes: float64(3), int64(35), object(43)
memory usage: 924.0+ KB

```


- This shows that there are null values present.
- Alley , PoolQC , MiscFeature and Fence have very less non-null values.
- Here target variable is SalePrice

```

In [5]: # target variable: SalePrice

sns.distplot(df['SalePrice'])
plt.show()

```



## 2. Data Cleaning

Let's now conduct some data cleaning steps.

```
In [6]: # duplicacy check  
df["Id"].is_unique
```

```
Out[6]: True
```

This means that no two Ids are same, hence we have all rows unique.

Let us check the percentage of null values.

```
In [7]: print(round(df.isnull().sum()/len(df.index)*100,2))
```

Id	0.00
MSSubClass	0.00
MSZoning	0.00
LotFrontage	17.74
LotArea	0.00
Street	0.00
Alley	93.77
LotShape	0.00
LandContour	0.00
Utilities	0.00
LotConfig	0.00
LandSlope	0.00
Neighborhood	0.00
Condition1	0.00
Condition2	0.00
BldgType	0.00
HouseStyle	0.00
OverallQual	0.00
OverallCond	0.00
YearBuilt	0.00
YearRemodAdd	0.00
RoofStyle	0.00
RoofMatl	0.00
Exterior1st	0.00
Exterior2nd	0.00
MasVnrType	0.55
MasVnrArea	0.55
ExterQual	0.00
ExterCond	0.00
Foundation	0.00
BsmtQual	2.53
BsmtCond	2.53
BsmtExposure	2.60
BsmtFinType1	2.53
BsmtFinSF1	0.00
BsmtFinType2	2.60
BsmtFinSF2	0.00
BsmtUnfSF	0.00
TotalBsmtSF	0.00
Heating	0.00
HeatingQC	0.00
CentralAir	0.00
Electrical	0.07
1stFlrSF	0.00
2ndFlrSF	0.00
LowQualFinSF	0.00
GrLivArea	0.00
BsmtFullBath	0.00
BsmtHalfBath	0.00
FullBath	0.00
HalfBath	0.00
BedroomAbvGr	0.00
KitchenAbvGr	0.00
KitchenQual	0.00
TotRmsAbvGrd	0.00
Functional	0.00
Fireplaces	0.00
FireplaceQu	47.26
GarageType	5.55
GarageYrBlt	5.55

```
GarageFinish      5.55
GarageCars        0.00
GarageArea        0.00
GarageQual        5.55
GarageCond        5.55
PavedDrive        0.00
WoodDeckSF        0.00
OpenPorchSF       0.00
EnclosedPorch     0.00
3SsnPorch         0.00
ScreenPorch       0.00
PoolArea          0.00
PoolQC           99.52
Fence             80.75
MiscFeature       96.30
MiscVal           0.00
MoSold            0.00
YrSold            0.00
SaleType          0.00
SaleCondition     0.00
SalePrice         0.00
dtype: float64
```

We have seen in the data description that there are few 'NA' values present. These are not null values rather they are 'not available' values. For e.g. NA in garage variables mean that there is no garage present in that house.

So we need to impute these wherever present so that it is not counted as null values.

We have seen that Alley, PoolQC, MiscFeature and Fence have more than 80% missing values. Even if we impute the NA with some other categorical value, that will become the dominating value. Since the percentage is so high, we can drop these columns.

```
In [8]: df = df.drop( columns = ['Alley', 'PoolQC', 'MiscFeature', 'Fence'])
df.info()
```

<class 'pandas.core.frame.DataFrame'>

RangeIndex: 1460 entries, 0 to 1459

Data columns (total 77 columns):

#	Column	Non-Null Count	Dtype
0	Id	1460 non-null	int64
1	MSSubClass	1460 non-null	int64
2	MSZoning	1460 non-null	object
3	LotFrontage	1201 non-null	float64
4	LotArea	1460 non-null	int64
5	Street	1460 non-null	object
6	LotShape	1460 non-null	object
7	LandContour	1460 non-null	object
8	Utilities	1460 non-null	object
9	LotConfig	1460 non-null	object
10	LandSlope	1460 non-null	object
11	Neighborhood	1460 non-null	object
12	Condition1	1460 non-null	object
13	Condition2	1460 non-null	object
14	BldgType	1460 non-null	object
15	HouseStyle	1460 non-null	object
16	OverallQual	1460 non-null	int64
17	OverallCond	1460 non-null	int64
18	YearBuilt	1460 non-null	int64
19	YearRemodAdd	1460 non-null	int64
20	RoofStyle	1460 non-null	object
21	RoofMatl	1460 non-null	object
22	Exterior1st	1460 non-null	object
23	Exterior2nd	1460 non-null	object
24	MasVnrType	1452 non-null	object
25	MasVnrArea	1452 non-null	float64
26	ExterQual	1460 non-null	object
27	ExterCond	1460 non-null	object
28	Foundation	1460 non-null	object
29	BsmtQual	1423 non-null	object
30	BsmtCond	1423 non-null	object
31	BsmtExposure	1422 non-null	object
32	BsmtFinType1	1423 non-null	object
33	BsmtFinSF1	1460 non-null	int64
34	BsmtFinType2	1422 non-null	object
35	BsmtFinSF2	1460 non-null	int64
36	BsmtUnfSF	1460 non-null	int64
37	TotalBsmtSF	1460 non-null	int64
38	Heating	1460 non-null	object
39	HeatingQC	1460 non-null	object
40	CentralAir	1460 non-null	object
41	Electrical	1459 non-null	object
42	1stFlrSF	1460 non-null	int64
43	2ndFlrSF	1460 non-null	int64
44	LowQualFinSF	1460 non-null	int64
45	GrLivArea	1460 non-null	int64
46	BsmtFullBath	1460 non-null	int64
47	BsmtHalfBath	1460 non-null	int64
48	FullBath	1460 non-null	int64
49	HalfBath	1460 non-null	int64
50	BedroomAbvGr	1460 non-null	int64
51	KitchenAbvGr	1460 non-null	int64
52	KitchenQual	1460 non-null	object
53	TotRmsAbvGrd	1460 non-null	int64
54	Functional	1460 non-null	object



```

55 Fireplaces      1460 non-null    int64
56 FireplaceQu     770 non-null    object
57 GarageType      1379 non-null    object
58 GarageYrBlt     1379 non-null    float64
59 GarageFinish    1379 non-null    object
60 GarageCars      1460 non-null    int64
61 GarageArea      1460 non-null    int64
62 GarageQual      1379 non-null    object
63 GarageCond      1379 non-null    object
64 PavedDrive      1460 non-null    object
65 WoodDeckSF      1460 non-null    int64
66 OpenPorchSF     1460 non-null    int64
67 EnclosedPorch   1460 non-null    int64
68 3SsnPorch       1460 non-null    int64
69 ScreenPorch     1460 non-null    int64
70 PoolArea        1460 non-null    int64
71 MiscVal         1460 non-null    int64
72 MoSold          1460 non-null    int64
73 YrSold          1460 non-null    int64
74 SaleType        1460 non-null    object
75 SaleCondition   1460 non-null    object
76 SalePrice       1460 non-null    int64
dtypes: float64(3), int64(35), object(39)
memory usage: 878.4+ KB

```

```
In [9]: df.shape
```

```
Out[9]: (1460, 77)
```

Now we have 77 columns. Lets start imputing the NA values now one by one.

## Garage

Lets talk about Garage related variables i.e. `GarageType` , `GarageFinish` , `GarageQual` and `GarageCond` .

```
In [10]: print('#NA in Garage Type :',df.GarageType.isnull().sum())
print('#NA in Garage Finish :',df.GarageFinish.isnull().sum())
print('#NA in Garage Qual :',df.GarageQual.isnull().sum())
print('#NA in Garage Cond :',df.GarageCond.isnull().sum())
```

```

#NA in Garage Type : 81
#NA in Garage Finish : 81
#NA in Garage Qual : 81
#NA in Garage Cond : 81

```

This means that 81 houses don't have garage. We can impute these with some other value like NoGarage.

```
In [11]: columns = ['GarageType', 'GarageFinish', 'GarageQual','GarageCond']
for col in columns:
    df[col].fillna('NoGarage',inplace=True)
```

```
In [12]: print('#NA in Garage Type :',df.GarageType.isnull().sum())
print('#NA in Garage Finish :',df.GarageFinish.isnull().sum())
print('#NA in Garage Qual :',df.GarageQual.isnull().sum())
print('#NA in Garage Cond :',df.GarageCond.isnull().sum())
```

```
#NA in Garage Type : 0
#NA in Garage Finish : 0
#NA in Garage Qual : 0
#NA in Garage Cond : 0
```

## Basement

Similarly we have to do the same thing with Basement.

Basement related variables are : `BsmtQual` , `BsmtCond` , `BsmtExposure` , `BsmtFinType1` , `BsmtFinType2`

```
In [13]: print('#NA in BsmtQual : ',df.BsmtQual.isnull().sum())
print('#NA in BsmtCond : ',df.BsmtCond.isnull().sum())
print('#NA in BsmtExposure : ',df.BsmtExposure.isnull().sum())
print('#NA in BsmtFinType1 : ',df.BsmtFinType1.isnull().sum())
print('#NA in BsmtFinType2 : ',df.BsmtFinType2.isnull().sum())

#NA in BsmtQual : 37
#NA in BsmtCond : 37
#NA in BsmtExposure : 38
#NA in BsmtFinType1 : 37
#NA in BsmtFinType2 : 38
```

```
In [14]: bsmt_columns = ['BsmtQual', 'BsmtCond', 'BsmtExposure', 'BsmtFinType1', 'BsmtFinType2']
for col in bsmt_columns:
    df[col].fillna('NoBasement',inplace=True)
```

## GarageYrBlt

Now, lets check `GarageYrBlt` as it has around 5.5% NA values.

```
In [15]: df.GarageYrBlt.describe()

Out[15]: count    1379.000000
mean      1978.506164
std        24.689725
min        1900.000000
25%        1961.000000
50%        1980.000000
75%        2002.000000
max        2010.000000
Name: GarageYrBlt, dtype: float64
```

```
In [16]: df.GarageYrBlt.isnull().sum()
```

```
Out[16]: 81
```

```
In [17]: df['GarageYrBlt'].fillna(df.GarageYrBlt.median(),inplace=True)
```

```
In [18]: df.GarageYrBlt.isnull().sum()
```

```
Out[18]: 0
```

## FireplaceQu

we can see `FireplaceQu` has 47% null values. So lets check the `value_counts` for that.

```
In [19]: df.FireplaceQu.value_counts()
```

```
Out[19]: Gd      380  
TA       313  
Fa        33  
Ex        24  
Po        20  
Name: FireplaceQu, dtype: int64
```

```
In [20]: df.FireplaceQu.isnull().sum()
```

```
Out[20]: 690
```

690 out of 1460 values are NA in this and NA means No fireplace according to the data description. So lets impute this.

```
In [21]: df['FireplaceQu'].fillna('NoFireplace',inplace=True)
```

```
In [22]: df.FireplaceQu.value_counts()
```

```
Out[22]: NoFireplace    690  
Gd      380  
TA       313  
Fa        33  
Ex        24  
Po        20  
Name: FireplaceQu, dtype: int64
```

## LotFrontage

Now, lets check `LotFrontage` as it has 17% NA values.

```
In [23]: df.LotFrontage.describe()
```

```
Out[23]: count      1201.000000  
mean         70.049958  
std          24.284752  
min          21.000000  
25%          59.000000  
50%          69.000000  
75%          80.000000  
max          313.000000  
Name: LotFrontage, dtype: float64
```

```
In [24]: df['LotFrontage'].isnull().sum()
```

```
Out[24]: 259
```

```
In [25]: df['LotFrontage'].fillna(df.LotFrontage.median(),inplace=True)
```

```
In [26]: df['LotFrontage'].isnull().sum()
```

```
Out[26]: 0
```

## MasVnrType

```
In [27]: df.MasVnrType.isnull().sum()
```

```
Out[27]: 8
```

```
In [28]: df.MasVnrType.value_counts()
```

```
Out[28]: None      864  
BrkFace   445  
Stone     128  
BrkCmn     15  
Name: MasVnrType, dtype: int64
```

```
In [29]: df.MasVnrType.fillna(df['MasVnrType'].mode()[0],inplace=True)
```

```
In [30]: df.MasVnrType.isnull().sum()
```

```
Out[30]: 0
```

## MasVnrArea

```
In [31]: print("#Null in MasVnrArea :",df.MasVnrArea.isnull().sum())
```

```
#Null in MasVnrArea : 8
```

```
In [32]: df['MasVnrArea'].fillna(df.MasVnrArea.median(),inplace=True)
```

```
In [33]: df.MasVnrArea.isnull().sum()
```

```
Out[33]: 0
```

## Electrical

```
In [34]: print("#Null in Electrical :",df.Electrical.isnull().sum())
```

```
#Null in Electrical : 1
```

```
In [35]: df.Electrical.value_counts()
```

```
Out[35]: SBrkr     1334  
FuseA       94  
FuseF       27  
FuseP        3  
Mix          1  
Name: Electrical, dtype: int64
```

```
In [36]: df.Electrical.fillna(df['Electrical'].mode()[0],inplace=True)
```

```
In [37]: df.Electrical.value_counts()
```

```
Out[37]: SBrkr      1335  
         FuseA       94  
         FuseF       27  
         FuseP        3  
         Mix         1  
         Name: Electrical, dtype: int64
```

```
In [38]: #Let us check the percentage of null values again  
  
         print(round(df.isnull().sum()/len(df.index)*100,2))
```

Id	0.0
MSSubClass	0.0
MSZoning	0.0
LotFrontage	0.0
LotArea	0.0
Street	0.0
LotShape	0.0
LandContour	0.0
Utilities	0.0
LotConfig	0.0
LandSlope	0.0
Neighborhood	0.0
Condition1	0.0
Condition2	0.0
BldgType	0.0
HouseStyle	0.0
OverallQual	0.0
OverallCond	0.0
YearBuilt	0.0
YearRemodAdd	0.0
RoofStyle	0.0
RoofMatl	0.0
Exterior1st	0.0
Exterior2nd	0.0
MasVnrType	0.0
MasVnrArea	0.0
ExterQual	0.0
ExterCond	0.0
Foundation	0.0
BsmtQual	0.0
BsmtCond	0.0
BsmtExposure	0.0
BsmtFinType1	0.0
BsmtFinSF1	0.0
BsmtFinType2	0.0
BsmtFinSF2	0.0
BsmtUnfSF	0.0
TotalBsmtSF	0.0
Heating	0.0
HeatingQC	0.0
CentralAir	0.0
Electrical	0.0
1stFlrSF	0.0
2ndFlrSF	0.0
LowQualFinSF	0.0
GrLivArea	0.0
BsmtFullBath	0.0
BsmtHalfBath	0.0
FullBath	0.0
HalfBath	0.0
BedroomAbvGr	0.0
KitchenAbvGr	0.0
KitchenQual	0.0
TotRmsAbvGrd	0.0
Functional	0.0
Fireplaces	0.0
FireplaceQu	0.0
GarageType	0.0
GarageYrBlt	0.0
GarageFinish	0.0

```

GarageCars      0.0
GarageArea      0.0
GarageQual      0.0
GarageCond      0.0
PavedDrive      0.0
WoodDeckSF      0.0
OpenPorchSF     0.0
EnclosedPorch   0.0
3SsnPorch       0.0
ScreenPorch     0.0
PoolArea        0.0
MiscVal         0.0
MoSold          0.0
YrSold          0.0
SaleType        0.0
SaleCondition    0.0
SalePrice       0.0
dtype: float64

```

Alright, so all the NA values are now replaced and the unimportant columns are deleted.  
We have a clean dataset now.

```

In [39]: # all numeric (float and int) variables in the dataset

df_numeric = df.select_dtypes(include=['float64', 'int64'])
df_numeric.head()

```

```

Out[39]:
   Id  MSSubClass  LotFrontage  LotArea  OverallQual  OverallCond  YearBuilt  YearRemodAdd  MasV
0   1           60         65.0    8450             7             5       2003         2003
1   2           20         80.0    9600             6             8       1976         1976
2   3           60         68.0   11250             7             5       2001         2002
3   4           70         60.0    9550             7             5       1915         1970
4   5           60         84.0   14260             8             5       2000         2000

```

5 rows × 38 columns

```

In [40]: # correlation matrix
cor = df_numeric.corr()
cor

```

Out[40]:

	<b>Id</b>	<b>MSSubClass</b>	<b>LotFrontage</b>	<b>LotArea</b>	<b>OverallQual</b>	<b>OverallCond</b>	<b>YearBuilt</b>
<b>Id</b>	1.000000	0.011156	-0.009921	-0.033226	-0.028365	0.012609	-0.012713
<b>MSSubClass</b>	0.011156	1.000000	-0.356718	-0.139781	0.032628	-0.059316	0.027850
<b>LotFrontage</b>	-0.009921	-0.356718	1.000000	0.304522	0.234812	-0.053281	0.116685
<b>LotArea</b>	-0.033226	-0.139781	0.304522	1.000000	0.105806	-0.005636	0.014228
<b>OverallQual</b>	-0.028365	0.032628	0.234812	0.105806	1.000000	-0.091932	0.572323
<b>OverallCond</b>	0.012609	-0.059316	-0.053281	-0.005636	-0.091932	1.000000	-0.375983
<b>YearBuilt</b>	-0.012713	0.027850	0.116685	0.014228	0.572323	-0.375983	1.000000
<b>YearRemodAdd</b>	-0.021998	0.040581	0.083348	0.013788	0.550684	0.073741	0.592855
<b>MasVnrArea</b>	-0.051071	0.023573	0.178469	0.103321	0.407252	-0.125694	0.311600
<b>BsmtFinSF1</b>	-0.005024	-0.069836	0.214367	0.214103	0.239666	-0.046231	0.249503
<b>BsmtFinSF2</b>	-0.005968	-0.065649	0.042463	0.111170	-0.059119	0.040229	-0.049107
<b>BsmtUnfSF</b>	-0.007940	-0.140759	0.124098	-0.002618	0.308159	-0.136841	0.149040
<b>TotalBsmtSF</b>	-0.015415	-0.238518	0.363472	0.260833	0.537808	-0.171098	0.391452
<b>1stFlrSF</b>	0.010496	-0.251758	0.413773	0.299475	0.476224	-0.144203	0.281986
<b>2ndFlrSF</b>	0.005590	0.307886	0.072388	0.050986	0.295493	0.028942	0.010308
<b>LowQualFinSF</b>	-0.044230	0.046474	0.037469	0.004779	-0.030429	0.025494	-0.183784
<b>GrLivArea</b>	0.008273	0.074853	0.368007	0.263116	0.593007	-0.079686	0.199010
<b>BsmtFullBath</b>	0.002289	0.003491	0.090343	0.158155	0.111098	-0.054942	0.187599
<b>BsmtHalfBath</b>	-0.020155	-0.002333	-0.006979	0.048046	-0.040150	0.117821	-0.038162
<b>FullBath</b>	0.005587	0.131608	0.180534	0.126031	0.550600	-0.194149	0.468271
<b>HalfBath</b>	0.006784	0.177354	0.047222	0.014259	0.273458	-0.060769	0.242656
<b>BedroomAbvGr</b>	0.037719	-0.023438	0.236840	0.119690	0.101676	0.012980	-0.070651
<b>KitchenAbvGr</b>	0.002951	0.281721	-0.004905	-0.017784	-0.183882	-0.087001	-0.174800
<b>TotRmsAbvGrd</b>	0.027239	0.040380	0.320518	0.190015	0.427452	-0.057583	0.095589
<b>Fireplaces</b>	-0.019772	-0.045569	0.233221	0.271364	0.396765	-0.023820	0.147716
<b>GarageYrBlt</b>	-0.000122	0.081396	0.062996	-0.025865	0.514231	-0.306276	0.777182
<b>GarageCars</b>	0.016570	-0.040110	0.269539	0.154871	0.600671	-0.185758	0.537850
<b>GarageArea</b>	0.017634	-0.098672	0.323511	0.180403	0.562022	-0.151521	0.478954
<b>WoodDeckSF</b>	-0.029643	-0.012579	0.075542	0.171698	0.238923	-0.003334	0.224880
<b>OpenPorchSF</b>	-0.000477	-0.006100	0.137014	0.084774	0.308819	-0.032589	0.188686
<b>EnclosedPorch</b>	0.002889	-0.012037	0.010287	-0.018340	-0.113937	0.070356	-0.387268
<b>3SsnPorch</b>	-0.046635	-0.043825	0.061945	0.020423	0.030371	0.025504	0.031355
<b>ScreenPorch</b>	0.001330	-0.026030	0.037655	0.043160	0.064886	0.054811	-0.050364




	Id	MSSubClass	LotFrontage	LotArea	OverallQual	OverallCond	YearBuilt
PoolArea	0.057044	0.008283	0.180819	0.077672	0.065166	-0.001985	0.004950
MiscVal	-0.006242	-0.007683	-0.000255	0.038068	-0.031406	0.068777	-0.034383
MoSold	0.021172	-0.013585	0.010451	0.001205	0.070815	-0.003511	0.012398
YrSold	0.000712	-0.021407	0.006380	-0.014261	-0.027347	0.043950	-0.013618
SalePrice	-0.021917	-0.084284	0.334771	0.263843	0.790982	-0.077856	0.522897

20 20 1

```
In [41]: # plotting correlations on a heatmap

# figure size
plt.figure(figsize=(30,20))

# heatmap
sns.heatmap(cor, cmap="YlGnBu", annot=True)
plt.show()
```



If we concentrate on the last row, we can see the following variables are highly correlated with target variable `SalePrice` :

It is positively correlated with following variables:

- OverallQual

- YearBuilt
- YearRemodAdd
- MasVnrArea
- TotalBsmtSF
- 1stFlrSF
- GrLivArea
- FullBath
- TotRmsAbvGrd
- GarageCars
- GarageArea

It is negatively correlated with following variables:

- Id
- MSSubClass
- OverallCond
- LowQualFinSF
- BsmtHalfBath
- KitchenAbvGr
- EnclosedPorch
- MiscVal
- YrSold

There is some multicollinearity present as well:

- 1stFlrSF and TotalBsmtSF
- GarageCars and GarageArea
- GrLivArea and TotRmsAbvGrd

and so on..

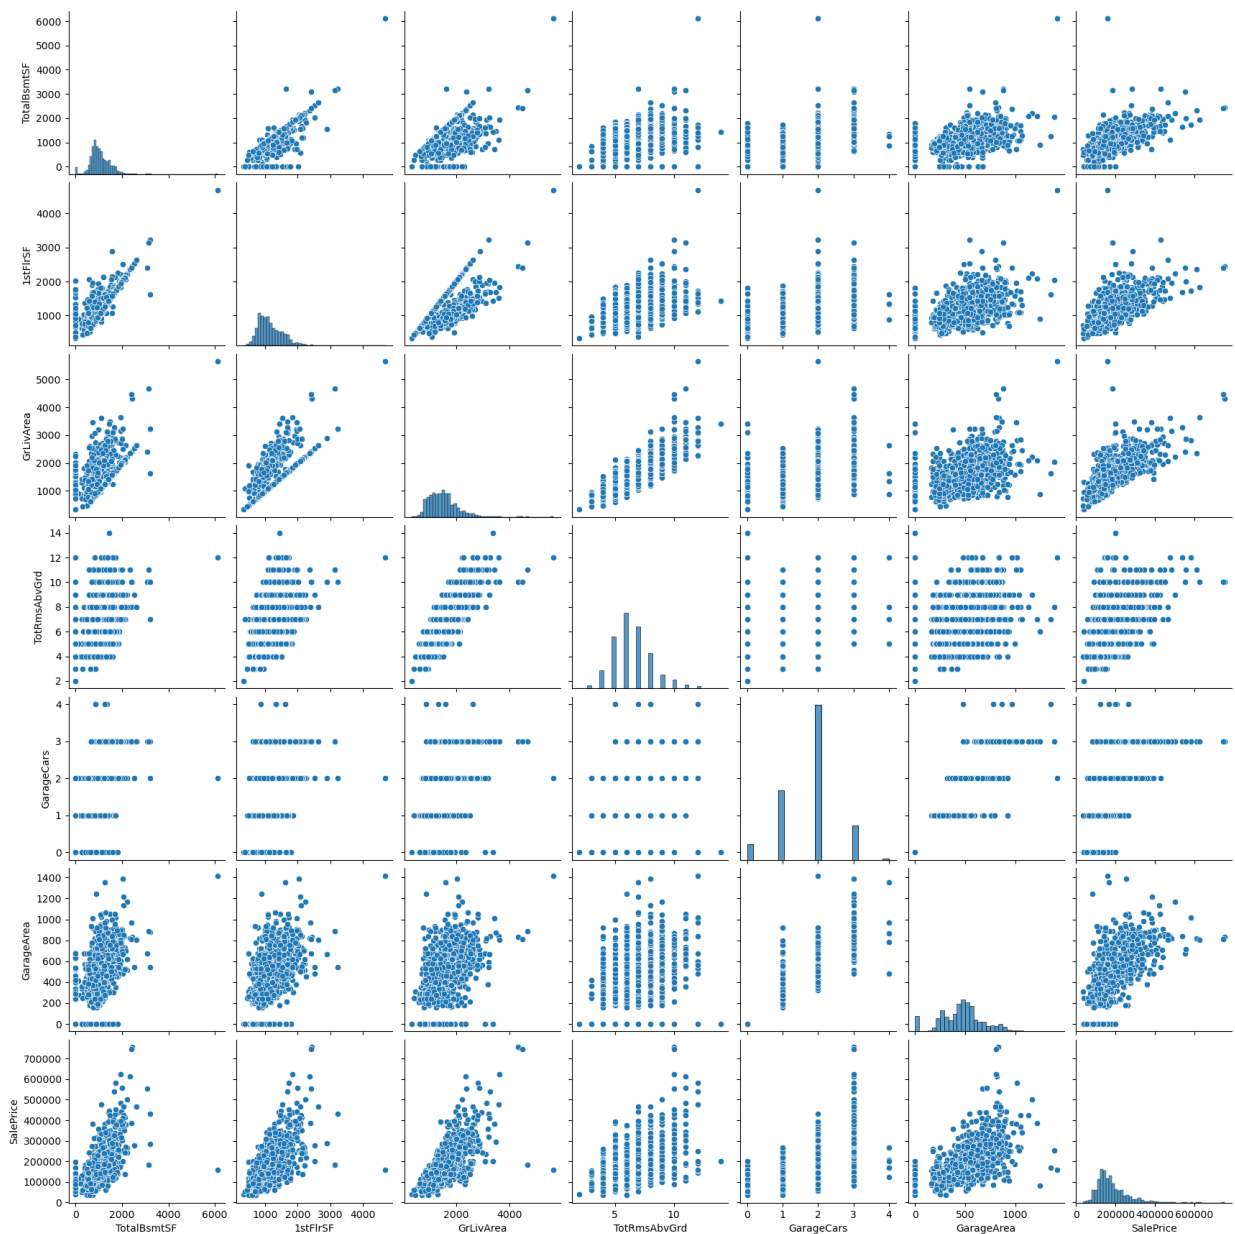
We will keep these points in mind while building the model.

since finding the multicollinearity here will be little difficult. Lets try to find top variables which has a correlation greater than 80%.

```
In [42]: cor = cor.abs()
top_cor_variables = np.where(cor>0.8)
top_cor_variables = [(cor.columns[x],cor.columns[y]) for x,y in zip(*top_cor_variables)]
print(top_cor_variables)
```

```
[('TotalBsmtSF', '1stFlrSF'), ('GrLivArea', 'TotRmsAbvGrd'), ('GarageCars', 'GarageArea')]
```

```
In [43]: # Lets do a pairplot to check the patterns
cols = ['TotalBsmtSF', '1stFlrSF', 'GrLivArea', 'TotRmsAbvGrd', 'GarageCars', 'GarageArea']
sns.pairplot(df[cols])
plt.show()
```



From pairplot, we can say `TotRmsAbvGrd` and `GarageCars` are not related to `SalePrice` and hence we can delete these.

```
In [44]: df = df.drop( columns = ['TotRmsAbvGrd', 'GarageArea'])
df.shape
```

```
Out[44]: (1460, 75)
```

## Derived Variables

Lets try to find the derived variables now.

One can be `AgeOfHouse` (age of the house) which can be formulated from `YrSold` minus `YearBuilt`.

```
In [45]: # create a new column AgeOfHouse
df['AgeOfHouse'] = df.YrSold - df.YearBuilt
```

```
# drop current year and year built columns
df.drop(['YrSold', 'YearBuilt'], axis=1, inplace=True)
```


We can drop the Id column as that is also not going to add any value to the model building.

```
In [46]: df = df.drop(['Id'], axis=1)
df.shape
```

```
Out[46]: (1460, 73)
```

```
In [47]: df['SalePrice'] = np.log1p(df['SalePrice'])
sns.distplot(df['SalePrice'])
```

```
Out[47]: <Axes: xlabel='SalePrice', ylabel='Density'>
```



## 3. Data Preparation

### Data Preparation

Let's now prepare the data and build the model.

```
In [48]: # split into X and y

# predictors in variable X
X = df.drop('SalePrice', axis=1)

# response variable in Y
y = df['SalePrice']
```

```
In [49]: # creating dummy variables for categorical variables
```

```
# subset all categorical variables
df_categorical = X.select_dtypes(include=['object'])
df_categorical.head()
```

```
Out[49]:
```

	MSZoning	Street	LotShape	LandContour	Utilities	LotConfig	LandSlope	Neighborhood	Condition
0	RL	Pave	Reg	Lvl	AllPub	Inside	Gtl	CollgCr	N
1	RL	Pave	Reg	Lvl	AllPub	FR2	Gtl	Veenker	F
2	RL	Pave	IR1	Lvl	AllPub	Inside	Gtl	CollgCr	N
3	RL	Pave	IR1	Lvl	AllPub	Corner	Gtl	Crawfor	N
4	RL	Pave	IR1	Lvl	AllPub	FR2	Gtl	NoRidge	N

5 rows × 39 columns

```
In [50]: # convert into dummies - one hot encoding
df_dummies = pd.get_dummies(df_categorical, drop_first=True)
df_dummies.head()
```

```
Out[50]:
```

	MSZoning_FV	MSZoning_RH	MSZoning_RL	MSZoning_RM	Street_Pave	LotShape_IR2	LotShape_IL
0	0	0	1	0	1	0	
1	0	0	1	0	1	0	
2	0	0	1	0	1	0	
3	0	0	1	0	1	0	
4	0	0	1	0	1	0	

5 rows × 210 columns

```
In [51]: # drop categorical variables
X = X.drop(list(df_categorical.columns), axis=1)
```

```
In [52]: # concat dummy variables with X
X = pd.concat([X, df_dummies], axis=1)
```

```
In [53]: # scaling the features - necessary before using Ridge or Lasso
from sklearn.preprocessing import scale

# storing column names in cols, since column names are lost after
# scaling (the df is converted to a numpy array)
cols = X.columns
X = pd.DataFrame(scale(X))
X.columns = cols
X.columns
```

```
Out[53]: Index(['MSSubClass', 'LotFrontage', 'LotArea', 'OverallQual', 'OverallCond',
        'YearRemodAdd', 'MasVnrArea', 'BsmtFinSF1', 'BsmtFinSF2', 'BsmtUnfSF',
        ...,
        'SaleType_ConLI', 'SaleType_ConLw', 'SaleType_New', 'SaleType_Oth',
        'SaleType_WD', 'SaleCondition_AdjLand', 'SaleCondition_Alloca',
        'SaleCondition_Family', 'SaleCondition_Normal',
        'SaleCondition_Partial'],
        dtype='object', length=243)
```

```
In [54]: # split into train and test
from sklearn.model_selection import train_test_split
X_train, X_test, y_train, y_test = train_test_split(X, y, train_size=0.7, test_size = 0.3)
```

## Scaling

```
In [55]: from sklearn.preprocessing import StandardScaler
scaler = StandardScaler()
```

Lets check the number of numerical and categorical values.

```
In [56]: # Checking number of numerical and categorical values

numerical_var = df.dtypes[df.dtypes != 'object'].index
print("Number of Numerical Variables : ", len(numerical_var))
print("Numerical variables : ", numerical_var)

categorical_var = df.dtypes[df.dtypes == 'object'].index
print("Number of Categorical Variables : " , len(categorical_var))
print("Categorical variables : ", categorical_var)
```

```
Number of Numerical Variables : 34
Numerical variables : Index(['MSSubClass', 'LotFrontage', 'LotArea', 'OverallQual',
        'OverallCond',
        'YearRemodAdd', 'MasVnrArea', 'BsmtFinSF1', 'BsmtFinSF2', 'BsmtUnfSF',
        'TotalBsmtSF', '1stFlrSF', '2ndFlrSF', 'LowQualFinSF', 'GrLivArea',
        'BsmtFullBath', 'BsmtHalfBath', 'FullBath', 'HalfBath', 'BedroomAbvGr',
        'KitchenAbvGr', 'Fireplaces', 'GarageYrBlt', 'GarageCars', 'WoodDeckSF',
        'OpenPorchSF', 'EnclosedPorch', '3SsnPorch', 'ScreenPorch', 'PoolArea',
        'MiscVal', 'MoSold', 'SalePrice', 'AgeOfHouse'],
        dtype='object')
Number of Categorical Variables : 39
Categorical variables : Index(['MSZoning', 'Street', 'LotShape', 'LandContour', 'Utilities',
        'LotConfig', 'LandSlope', 'Neighborhood', 'Condition1', 'Condition2',
        'BldgType', 'HouseStyle', 'RoofStyle', 'RoofMatl', 'Exterior1st',
        'Exterior2nd', 'MasVnrType', 'ExterQual', 'ExterCond', 'Foundation',
        'BsmtQual', 'BsmtCond', 'BsmtExposure', 'BsmtFinType1', 'BsmtFinType2',
        'Heating', 'HeatingQC', 'CentralAir', 'Electrical', 'KitchenQual',
        'Functional', 'FireplaceQu', 'GarageType', 'GarageFinish', 'GarageQual',
        'GarageCond', 'PavedDrive', 'SaleType', 'SaleCondition'],
        dtype='object')
```

So we have 34 numerical values and 39 categorical values.

```
In [57]: # removing target variable

numerical_var = list(numerical_var)
```

```
numerical_var.remove('SalePrice')
```

```
# fit transform on train set
```

```
X_train[numerical_var] = scaler.fit_transform(X_train[numerical_var])
```

```
X_train.head()
```

```
Out[57]:
```

	MSSubClass	LotFrontage	LotArea	OverallQual	OverallCond	YearRemodAdd	MasVnrArea	B
210	-0.657071	-0.115302	-0.473765	-0.779861	0.383154	-1.694350	-0.558025	
318	0.035976	0.926898	-0.056845	0.649651	-0.533005	0.390956	0.809137	
239	-0.195040	-0.794998	-0.169324	-0.065105	-1.449164	-1.694350	-0.558025	
986	-0.195040	-0.477806	-0.502297	-0.065105	2.215472	0.875911	-0.558025	
1416	3.039179	-0.432493	0.082905	-1.494617	0.383154	-1.694350	-0.558025	

5 rows × 243 columns

```
In [58]: # transform on test set
```

```
X_test[numerical_var] = scaler.transform(X_test[numerical_var])
```

```
X_test.head()
```

```
Out[58]:
```

	MSSubClass	LotFrontage	LotArea	OverallQual	OverallCond	YearRemodAdd	MasVnrArea	B
1436	-0.888086	-0.432493	-0.144189	-1.494617	0.383154	-0.675945	-0.558025	
57	0.035976	0.881585	0.112505	0.649651	-0.533005	0.924407	-0.558025	
780	-0.888086	-0.296554	-0.253368	0.649651	-0.533005	0.536443	-0.355087	
382	0.035976	0.428455	-0.120412	0.649651	-0.533005	1.021398	-0.558025	
1170	0.498007	0.292515	-0.058786	-0.065105	0.383154	-0.384972	-0.558025	

5 rows × 243 columns

## 4. Model Building and Evaluation

### Linear Regression

Let's now try predicting car prices, a dataset using linear regression.

```
In [59]: # Instantiate
```

```
lm = LinearRegression()
```

```
# Fit a Line
```

```
lm.fit(X_train, y_train)
```

Out[59]:

▼ LinearRegression

LinearRegression()

In [60]:

```
# Print the coefficients and intercept  
print(lm.intercept_)  
print(lm.coef_)
```



-3179715.3690656256

[-1.52060786e-02	7.81767943e-03	3.47357315e-02	5.58230959e-02
4.02892434e-02	2.36918099e-02	1.16774065e-05	1.74605452e+09
6.41614605e+08	1.68495629e+09	-1.72427926e+09	-1.06235997e+08
-1.16189582e+08	-1.28675882e+07	1.42005328e+08	1.23034865e-02
6.69278204e-04	3.19680572e-03	2.25646794e-03	7.33169913e-03
-1.43100321e-02	-5.29515743e-03	4.39584255e-03	2.91034728e-02
1.01996064e-02	4.07002866e-03	6.85892999e-03	3.81037593e-03
8.09580274e-03	2.90453248e-03	1.02119148e-03	7.19979405e-04
-6.23948425e-02	9.27400291e-02	4.63690683e-02	1.79391026e-01
1.35035396e-01	5.74129447e-03	2.45153904e-03	-1.29853934e-03
2.47024000e-03	1.68631598e-03	-6.87431544e-04	7.21601397e-03
-2.04801746e-03	7.16167688e-03	-7.26862694e-03	-3.52799892e-04
-2.84837186e-03	4.58277017e-03	-2.06102878e-02	2.84323451e-03
-7.55995512e-04	1.99696720e-02	1.05839297e-02	7.71742314e-03
3.35295796e-02	-3.74168903e-03	3.47119570e-03	1.11695826e-02
-1.19463205e-02	-4.80495393e-04	1.24228746e-02	1.24412775e-03
1.60464644e-03	7.31796026e-03	1.67711377e-02	1.90507174e-02
1.28956288e-02	3.16837430e-03	6.47991896e-03	1.99429542e-02
1.72668733e-02	2.36503035e-03	8.44805688e-03	1.46470964e-02
3.49997878e-02	7.45096058e-03	8.13373178e-03	-3.75039876e-04
8.26363266e-03	1.94260478e-03	3.38380039e-03	3.80449276e-03
4.10519028e-03	5.93052804e-03	-6.37972802e-02	-1.48158595e-02
-1.69990957e-03	5.28885424e-03	-1.72934122e-03	-2.75892019e-03
-3.91066074e-03	5.25647402e-03	1.63689256e-04	-2.36412138e-02
-4.80136462e-03	2.09257007e-03	-7.57655501e-03	-3.33866477e-03
-2.55093724e-03	-2.59336233e-02	-5.29303402e-03	-2.47038603e-02
2.67957896e-03	1.62563771e-02	3.80282253e-01	8.54303185e-02
8.12032670e-02	7.75805861e-02	2.53355794e-01	1.61158755e-01
1.87016731e-01	-6.17242070e+07	-6.00449927e-03	2.97097862e-02
1.19180469e+08	-7.54803419e-04	2.79184878e-02	4.46081907e-03
3.10039222e-02	2.01646388e-02	5.56873903e-03	1.10934526e-02
4.52892110e-02	1.21361613e-02	1.39368773e-02	1.06836162e+08
5.82873821e-04	-8.31219554e-03	-1.19180469e+08	1.79044306e-02
-8.57262313e-03	-4.56400216e-03	-4.34249640e-03	-4.23457045e+04
-6.07657433e-03	-5.11620939e-03	-4.43309546e-04	-5.13602793e-03
2.16184556e-03	-7.15029240e-03	-4.36335802e-04	-5.33425063e-03
2.16078758e-03	2.38744915e-03	6.30659610e-03	5.91017678e-03
-1.28225982e-02	-2.18351483e-02	-3.91369914e+02	-1.70574207e-02
8.29645433e-03	1.34671312e-02	-6.68102503e-03	4.12029400e-03
-2.88008247e-03	-4.85509634e-04	-2.49008536e-02	6.88627262e+08
-2.18780041e-02	5.01921773e-03	-3.44313490e+08	1.15151517e-02
7.91320205e-03	1.08566061e-02	3.40916216e-04	-4.64257598e-03
-1.01690292e-02	-2.14829296e-03	5.58049977e-03	-5.98092750e-03
-3.44313771e+08	-7.01314211e-03	-7.60790706e-03	-1.22363567e-02
-5.22464514e-04	-7.25694560e-03	-2.61207223e-02	-6.88020885e-03
-6.99383020e-03	2.97570564e-02	2.57173944e-02	2.83662975e-03
5.84350526e-03	1.51006877e-02	-1.66263897e-04	-5.12066111e-03
-3.97010893e-03	-1.10913254e-02	1.03862528e-02	4.69550490e-03
-9.72416252e-04	5.01170020e-01	-1.56293064e-03	-8.28453153e-03
-3.37825902e-02	-2.87038162e-02	-7.74049014e-03	8.90510064e-03
1.39473975e-02	-5.04982471e-03	-1.08713210e-02	2.74178833e-02
2.42260098e-03	1.93149894e-02	1.68405473e-03	1.53884292e-03
1.73361301e-02	-6.00593537e-03	-4.98408824e-03	-4.32330370e-03
2.70904601e-03	-1.19801015e-02	1.34131180e+09	1.20140157e+09
-5.86435199e-04	-1.51333213e-03	2.03714644e+08	1.11335188e+08
-1.27135668e+09	5.17338365e+07	3.45842820e+08	-1.74753442e+08
-8.94209427e+07	-1.27135668e+09	-7.89161901e+07	-3.29843891e+08
4.55648452e-03	5.74502721e-03	2.77770311e-03	3.64865363e-03
2.03541517e-02	-2.71950383e-04	8.53111967e-04	7.57951554e+08

```
2.20305473e-03  2.27481127e-04  5.86719252e-03  8.00196826e-03
4.26387787e-03  1.93721056e-02  -7.66353446e+08]
```

```
In [61]: y_pred_train = lm.predict(X_train)
y_pred_test = lm.predict(X_test)

metric = []
r2_train_lr = r2_score(y_train, y_pred_train)
print("R2 train:", r2_train_lr)
metric.append(r2_train_lr)

r2_test_lr = r2_score(y_test, y_pred_test)
print("R2 test:", r2_test_lr)
metric.append(r2_test_lr)

rss1_lr = np.sum(np.square(y_train - y_pred_train))
print("RSS train:", rss1_lr)
metric.append(rss1_lr)

rss2_lr = np.sum(np.square(y_test - y_pred_test))
print("RSS test:", rss2_lr)
metric.append(rss2_lr)

mse_train_lr = mean_squared_error(y_train, y_pred_train)
print("MSE train:", mse_train_lr)
metric.append(mse_train_lr)

print("RMSE train:", mse_train_lr**0.5)
metric.append(mse_train_lr**0.5)

mse_test_lr = mean_squared_error(y_test, y_pred_test)
print("MSE test:", mse_test_lr)
metric.append(mse_test_lr)

print("RMSE test:", mse_test_lr**0.5)
metric.append(mse_test_lr**0.5)
```

```
R2 train: 0.9581906935195057
R2 test: -4.8487018426528915e+17
RSS train: 6.709983268507258
RSS test: 3.4943434309044646e+19
MSE train: 0.006571971859458626
RMSE train: 0.08106769923624714
MSE test: 7.977953038594669e+16
RMSE test: 282452704.68867296
```

This shows that r2 value is very poor.

## Ridge and Lasso Regression

Let's now try predicting sale prices, a dataset used in simple linear regression, to perform ridge and lasso regression.

## Ridge Regression

```
In [62]: # List of alphas to tune - if value too high it will lead to underfitting, if it is too low
# it will not handle the overfitting
params = {'alpha': [0.0001, 0.001, 0.01, 0.05, 0.1,
                    0.2, 0.3, 0.4, 0.5, 0.6, 0.7, 0.8, 0.9, 1.0, 2.0, 3.0,
                    4.0, 5.0, 6.0, 7.0, 8.0, 9.0, 10.0, 20, 50, 100, 500, 1000 ]}

ridge = Ridge()

# cross validation
folds = 5
model_cv = GridSearchCV(estimator = ridge,
                        param_grid = params,
                        scoring= 'neg_mean_absolute_error',
                        cv = folds,
                        return_train_score=True,
                        verbose = 1)
model_cv.fit(X_train, y_train)
```

Fitting 5 folds for each of 28 candidates, totalling 140 fits

```
Out[62]: > GridSearchCV
> estimator: Ridge
    > Ridge
```

```
In [63]: # Printing the best hyperparameter alpha
print(model_cv.best_params_)

{'alpha': 4.0}
```

```
In [64]: #Fitting Ridge model for alpha = 4 and printing coefficients which have been penalised
alpha = 4
ridge = Ridge(alpha=alpha)

ridge.fit(X_train, y_train)
print(ridge.coef_)
```

[-1.27927128e-02 2.78527674e-03 3.30398298e-02 5.92514955e-02  
3.86146480e-02 2.33039749e-02 -4.30905825e-04 2.60513509e-02  
8.11483711e-03 7.07988597e-03 3.63183486e-02 4.78749433e-02  
4.88623875e-02 -1.15812446e-04 7.57848502e-02 1.53394349e-02  
8.84307927e-04 7.55386699e-03 4.80266940e-03 1.00108484e-02  
-1.43944220e-02 -7.10149128e-03 1.40926074e-04 3.33796478e-02  
1.04726614e-02 3.64937943e-03 7.35756683e-03 4.73971439e-03  
8.70281528e-03 -2.00889059e-03 9.39217954e-04 4.37557636e-04  
-5.19053325e-02 6.86833696e-02 3.52822521e-02 1.34376493e-01  
1.01119572e-01 6.31910579e-03 2.67457495e-03 -3.90556373e-03  
2.32886074e-03 3.42691741e-03 9.49663119e-04 9.56074962e-03  
-2.48220657e-03 7.14227116e-03 -8.25964535e-03 -1.44238795e-03  
-3.43119416e-03 4.25619570e-03 -1.78083531e-02 1.32599189e-03  
-3.41865522e-03 1.23282186e-02 9.70435682e-03 4.37942814e-03  
3.01039053e-02 -1.00679120e-02 8.12322744e-04 1.51314210e-03  
-1.58194298e-02 -3.32651137e-03 6.34618250e-03 -4.70521870e-04  
-1.29333615e-03 8.11196781e-03 1.83695162e-02 7.90036798e-03  
9.21391958e-03 -6.83962949e-04 4.32838551e-03 1.97509864e-02  
1.62143051e-02 1.08563656e-03 7.57003747e-03 1.18934061e-02  
3.36456519e-02 7.43100021e-03 7.69513170e-03 -7.67772984e-04  
8.88044438e-03 1.66769429e-03 4.60007574e-03 1.96298349e-03  
2.32248062e-03 6.85467887e-03 -5.89694705e-02 -1.45274101e-02  
-1.95737502e-03 4.21150475e-03 -2.26981378e-03 -1.55914707e-03  
-7.44564784e-03 -4.68726376e-04 3.79915853e-04 -1.59796739e-02  
-4.86381915e-03 3.86995247e-04 -9.88107060e-03 -2.65949706e-03  
-3.46400189e-03 -2.25839560e-02 -4.19005854e-03 -1.98229642e-02  
3.31766585e-03 1.58543429e-02 2.75774300e-01 6.27153463e-02  
5.96421637e-02 5.67798788e-02 1.83106336e-01 1.13698604e-01  
1.37017522e-01 3.51909574e-04 -7.23862534e-03 2.19742283e-02  
-6.06695487e-05 -4.66755531e-03 1.37685404e-02 2.61921485e-03  
1.64779918e-02 1.12793965e-02 4.69200997e-03 5.12652023e-03  
2.95841788e-02 8.56945906e-04 8.90511151e-03 2.03314487e-04  
2.45785160e-03 -3.52509160e-03 -6.06695487e-05 2.16688879e-02  
2.42474762e-03 -2.07221157e-04 7.38480611e-03 0.00000000e+00  
2.68390787e-03 -3.97172789e-03 1.15693857e-03 7.03543754e-03  
1.07246569e-02 -3.84533640e-03 7.27065795e-04 -3.35405263e-03  
1.74033502e-03 1.11958625e-03 1.03099426e-02 8.97987335e-03  
-7.42280501e-03 -1.05490974e-02 0.00000000e+00 -5.50621464e-03  
1.11232372e-02 1.55370398e-02 -4.97969535e-03 3.50827150e-03  
-1.96434631e-03 -2.08345295e-03 -2.67497341e-02 6.39580988e-03  
-2.56234073e-02 5.56877957e-03 6.39580988e-03 9.91830835e-03  
8.01717380e-03 1.18187369e-02 -9.93609486e-04 -5.97797185e-03  
-8.69854131e-03 -3.19861822e-03 6.04602252e-03 -6.09994043e-03  
6.39580988e-03 -7.79228842e-03 -1.22730537e-02 -1.22291941e-02  
-4.94931726e-04 -5.96763627e-03 -1.64359786e-02 -6.64512014e-03  
-6.60979171e-03 2.04095047e-02 2.01105109e-02 -9.52136612e-04  
2.31748424e-03 1.11712204e-02 3.42553151e-04 -5.39431213e-03  
-3.68425696e-03 -1.17202750e-02 1.27117554e-02 5.25880737e-03  
-4.30016112e-04 0.00000000e+00 -1.31815784e-03 -8.33555538e-03  
-3.36135424e-02 -2.94196746e-02 -8.72815480e-03 6.47554518e-03  
1.04027767e-02 -4.25124545e-03 -1.02631377e-02 2.21061626e-02  
1.22760621e-03 1.26475662e-02 -8.92483192e-03 -7.68371652e-04  
1.11286072e-02 3.13546690e-03 -2.20994037e-03 -3.75484662e-04  
3.85619714e-03 -3.20336901e-03 -4.35771579e-03 -4.35771579e-03  
-1.07179440e-03 -3.72178155e-03 -1.60977826e-02 2.89412198e-03  
-4.35771579e-03 -2.68438550e-03 -7.00599676e-03 -9.90589003e-03  
-3.67482235e-03 -4.35771579e-03 4.10348697e-03 -1.14694281e-02  
5.31747621e-03 5.96697003e-03 3.23026477e-03 4.10810637e-03  
1.95007371e-02 1.96927340e-04 1.08746589e-03 1.38319565e-02

```
2.51181962e-03  2.17392111e-03  4.74201485e-03  6.61277205e-03
4.57924472e-03  1.97939383e-02  1.36803103e-02]
```

In [65]: *# Lets calculate some metrics such as R2 score, RSS and RMSE*

```
y_pred_train = ridge.predict(X_train)
y_pred_test = ridge.predict(X_test)

metric2 = []
r2_train_lr = r2_score(y_train, y_pred_train)
print("R square train :",r2_train_lr)
metric2.append(r2_train_lr)

r2_test_lr = r2_score(y_test, y_pred_test)
print("R square test :",r2_test_lr)
metric2.append(r2_test_lr)

rss1_lr = np.sum(np.square(y_train - y_pred_train))
print("Rss train :",rss1_lr)
metric2.append(rss1_lr)

rss2_lr = np.sum(np.square(y_test - y_pred_test))
print("Rss test :",rss2_lr)
metric2.append(rss2_lr)

mse_train_lr = mean_squared_error(y_train, y_pred_train)
print("MSE train:",mse_train_lr)
metric2.append(mse_train_lr)

print("RMSE train:",mse_train_lr**0.5)
metric2.append(mse_train_lr**0.5)


mse_test_lr = mean_squared_error(y_test, y_pred_test)
print("MSE test :",mse_test_lr)
metric2.append(mse_test_lr)

print("RMSE test :",mse_test_lr**0.5)
metric2.append(mse_test_lr**0.5)
```

```
R square train : 0.9563668026337852
R square test : 0.8692068548139035
Rss train : 7.002699851416305
Rss test : 9.425949099776162
MSE train: 0.0068586678270482915
RMSE train: 0.08281707448979522
MSE test : 0.021520431734648772
RMSE test : 0.14669843807842253
```

In [66]: *#plotting*

```
model_cv_results = pd.DataFrame(model_cv.cv_results_)
plt.plot(model_cv_results['param_alpha'], model_cv_results['mean_train_score'])
plt.plot(model_cv_results['param_alpha'], model_cv_results['mean_test_score'])
plt.xlabel('alpha')
plt.xscale('log')
plt.ylabel('R2 value')
plt.legend(['train score', 'test score'])
plt.show()
```



As alpha increases:

- train error decreases
- test error is first constant and then decreases

## Lasso

```
In [67]: lasso = Lasso()

# cross validation
model_cv = GridSearchCV(estimator = lasso,
                        param_grid = params,
                        scoring= 'neg_mean_absolute_error',
                        cv = folds,
                        return_train_score=True,
                        verbose = 1)

model_cv.fit(X_train, y_train)
```

Fitting 5 folds for each of 28 candidates, totalling 140 fits

```
Out[67]: > GridSearchCV
> estimator: Lasso
```

```
> Lasso
```

```
In [68]: # Printing the best hyperparameter alpha
print(model_cv.best_params_)

{'alpha': 0.001}
```

```
In [69]: #Fitting Lasso model for alpha = 0.001 and printing coefficients which have been penal

alpha =0.001

lasso = Lasso(alpha=alpha)

lasso.fit(X_train, y_train)
```

```
Out[69]: ▾      Lasso
          Lasso(alpha=0.001)
```

```
In [70]: lasso.coef_
```

```
Out[70]: array([-1.61262086e-02,  2.08610968e-03,  2.41078698e-02,  6.62689614e-02,
  4.02439898e-02,  2.38003769e-02,  4.56033324e-05,  1.53927672e-02,
  8.55770266e-04, -0.00000000e+00,  3.79414235e-02,  2.40009372e-03,
  0.00000000e+00, -4.05746284e-03,  1.42040081e-01,  2.00867167e-02,
  0.00000000e+00,  6.13439927e-03,  4.09156761e-03,  4.62429597e-03,
 -1.17834084e-02,  0.00000000e+00,  0.00000000e+00,  3.46271101e-02,
  9.62235868e-03,  3.82877612e-03,  3.75925631e-03,  3.80934179e-03,
  7.12251361e-03, -4.79527133e-03, -0.00000000e+00, -0.00000000e+00,
 -5.29846777e-02,  3.52764960e-02,  1.88030062e-02,  7.44493007e-02,
  4.65672455e-02,  4.75788842e-03,  9.19110007e-04, -4.15465631e-03,
  0.00000000e+00,  0.00000000e+00,  0.00000000e+00,  3.70264764e-03,
 -2.15137733e-03,  7.52648211e-03, -4.45221578e-03, -3.11210762e-05,
 -5.96633914e-05,  1.35041959e-03, -8.20107100e-03,  0.00000000e+00,
 -3.55135636e-03,  6.24496247e-03,  9.28037160e-03,  0.00000000e+00,
  2.57092005e-02, -1.24746151e-02, -1.41143816e-03, -4.75963942e-03,
 -1.43563742e-02, -3.73970638e-03,  0.00000000e+00, -0.00000000e+00,
 -1.19258825e-03,  5.29331181e-03,  1.65398550e-02, -0.00000000e+00,
  5.15552926e-03, -1.69266478e-03,  6.07806334e-05,  1.80769194e-02,
  1.21485021e-02,  0.00000000e+00,  5.16149744e-03,  1.21791733e-03,
  2.03958084e-02,  3.02306083e-03,  2.88914232e-03, -3.99748206e-03,
  3.50187352e-03,  0.00000000e+00,  2.18511879e-03, -0.00000000e+00,
  9.91805972e-04,  5.14214865e-03, -5.77445259e-02, -6.11286847e-03,
 -1.45136196e-03,  1.91662622e-03,  0.00000000e+00, -0.00000000e+00,
 -8.91329620e-03, -2.35655659e-04,  9.84559746e-04, -4.06307132e-03,
 -3.81688409e-03,  0.00000000e+00, -0.00000000e+00, -0.00000000e+00,
  0.00000000e+00, -1.72762887e-03,  0.00000000e+00,  0.00000000e+00,
  4.57605452e-03,  7.73313960e-03,  2.42022528e-01,  5.20981205e-02,
  5.06481567e-02,  4.86084044e-02,  1.62587690e-01,  1.01794979e-01,
  1.21258260e-01, -0.00000000e+00, -7.72213839e-03,  1.21181421e-02,
 -1.81081156e-04,  0.00000000e+00, -1.80583138e-04,  3.55284433e-04,
  3.90653504e-03,  0.00000000e+00,  2.53523041e-04, -0.00000000e+00,
  1.06829526e-02, -3.69116002e-03,  2.87349844e-03, -0.00000000e+00,
  0.00000000e+00, -0.00000000e+00, -1.60628598e-05,  6.27686446e-03,
 -0.00000000e+00, -0.00000000e+00,  1.03941374e-03,  0.00000000e+00,
  0.00000000e+00, -0.00000000e+00, -1.04864690e-03,  0.00000000e+00,
  0.00000000e+00, -5.51140228e-03,  0.00000000e+00, -1.85423400e-03,
  2.88313141e-04, -1.79938406e-03,  1.78006929e-03, -3.81036264e-04,
 -2.24569740e-03, -0.00000000e+00,  0.00000000e+00,  4.37783219e-03,
  0.00000000e+00,  8.05034206e-03, -0.00000000e+00,  2.88123779e-03,
 -1.46899754e-03, -0.00000000e+00, -1.68610353e-02, -0.00000000e+00,
 -1.53277648e-02,  2.82542416e-03, -0.00000000e+00,  1.49444147e-03,
  5.42776782e-03,  1.48150976e-02, -0.00000000e+00, -4.85922669e-03,
 -6.83872767e-04, -0.00000000e+00,  7.93590656e-03, -1.04123296e-03,
 -0.00000000e+00, -4.17949632e-03, -8.47910269e-03, -7.05864663e-03,
  2.63415570e-03, -0.00000000e+00, -3.59514846e-03, -0.00000000e+00,
  0.00000000e+00,  0.00000000e+00,  4.35072819e-03, -6.44640908e-03,
 -2.52886597e-03,  1.32830375e-03,  0.00000000e+00, -3.21196360e-03,
 -5.92417825e-04, -9.00288106e-03,  1.57062901e-02,  3.24585104e-03,
 -1.07325705e-03,  0.00000000e+00,  0.00000000e+00, -3.11143648e-03,
 -1.91590424e-02, -1.51467692e-02, -9.75303722e-03,  0.00000000e+00,
  2.11615461e-03, -4.07795769e-03, -9.44256311e-03,  9.88706328e-03,
 -0.00000000e+00,  3.62237616e-03, -1.50910580e-02, -2.10682569e-03,
  0.00000000e+00,  4.67256764e-03, -1.22416573e-03,  0.00000000e+00,
  0.00000000e+00, -3.42428721e-04, -0.00000000e+00, -0.00000000e+00,
 -0.00000000e+00, -2.29807458e-03, -7.90308648e-03,  3.36753764e-03,
 -0.00000000e+00,  0.00000000e+00,  0.00000000e+00, -3.30761854e-03,
 -0.00000000e+00, -0.00000000e+00,  1.91484294e-03,  9.24963532e-04,
  4.33187613e-03,  4.80374297e-03,  2.69014042e-03,  2.28814145e-03,
  1.54591058e-02, -0.00000000e+00,  1.51238024e-04,  2.44847726e-02,
```



```
7.64416709e-04, 0.00000000e+00, 3.55156311e-03, 1.20613717e-03,  
2.45725324e-03, 1.79753584e-02, 5.20037031e-05])
```


```
In [71]: # Lets calculate some metrics such as R2 score, RSS and RMSE
```

```
y_pred_train = lasso.predict(X_train)  
y_pred_test = lasso.predict(X_test)  
  
metric3 = []  
r2_train_lr = r2_score(y_train, y_pred_train)  
print("R square train :",r2_train_lr)  
metric3.append(r2_train_lr)  
  
r2_test_lr = r2_score(y_test, y_pred_test)  
print("R square test :",r2_test_lr)  
metric3.append(r2_test_lr)  
  
rss1_lr = np.sum(np.square(y_train - y_pred_train))  
print("Rss train :",rss1_lr)  
metric3.append(rss1_lr)  
  
rss2_lr = np.sum(np.square(y_test - y_pred_test))  
print("Rss test :",rss2_lr)  
metric3.append(rss2_lr)  
  
mse_train_lr = mean_squared_error(y_train, y_pred_train)  
print("MSE train:",mse_train_lr)  
metric3.append(mse_train_lr)  
  
print("RMSE train:",mse_train_lr**0.5)  
metric3.append(mse_train_lr**0.5)  
  
mse_test_lr = mean_squared_error(y_test, y_pred_test)  
print("MSE test :",mse_test_lr)  
metric3.append(mse_test_lr)  
  
print("RMSE test :",mse_test_lr**0.5)  
metric3.append(mse_test_lr**0.5)
```

```
R square train : 0.9505553591166847  
R square test : 0.8724446319321947  
Rss train : 7.9353794877985155  
Rss test : 9.192610248034278  
MSE train: 0.007772164042897664  
RMSE train: 0.0881598777386724  
MSE test : 0.020987694630215246  
RMSE test : 0.14487130368094037
```

```
In [72]: model_cv_results2 = pd.DataFrame(model_cv.cv_results_)
```

```
#plotting  
plt.plot(model_cv_results2['param_alpha'], model_cv_results2['mean_train_score'])  
plt.plot(model_cv_results2['param_alpha'], model_cv_results2['mean_test_score'])  
plt.xlabel('alpha')  
plt.xscale('log')  
plt.ylabel('R2 value')  
plt.legend(['train score', 'test score'])  
plt.show()
```



With increase in alpha value, both train and test error decreases.

```
In [73]: # Creating a table which contain all the metrics

lr_table = {'Metric': ['R2 Score (Train)', 'R2 Score (Test)', 'RSS (Train)', 'RSS (Test)'],
            'Linear Regression': metric }

lr_metric = pd.DataFrame(lr_table, columns = ['Metric', 'Linear Regression'] )

rg_metric = pd.Series(metric2, name = 'Ridge Regression')
ls_metric = pd.Series(metric3, name = 'Lasso Regression')

final_metric = pd.concat([lr_metric, rg_metric, ls_metric], axis = 1)

final_metric
```

Out[73]:

	Metric	Linear Regression	Ridge Regression	Lasso Regression
0	R2 Score (Train)	9.581907e-01	0.956367	0.950555
1	R2 Score (Test)	-4.848702e+17	0.869207	0.872445
2	RSS (Train)	6.709983e+00	7.002700	7.935379
3	RSS (Test)	3.494343e+19	9.425949	9.192610
4	MSE (Train)	6.571972e-03	0.006859	0.007772
5	MSE (Test)	8.106770e-02	0.082817	0.088160
6	RMSE (Train)	7.977953e+16	0.021520	0.020988
7	RMSE (Test)	2.824527e+08	0.146698	0.144871

Lets observe the changes in the coefficients after regularization

```
In [74]: betas = pd.DataFrame(index=X.columns)
```

```
In [75]: betas.rows = X.columns
```

```
In [76]: betas['Linear'] = lm.coef_  
betas['Ridge'] = ridge.coef_  
betas['Lasso'] = lasso.coef_
```

```
In [77]: pd.set_option('display.max_rows', None)  
betas.head(80)
```

Out[77]:

	Linear	Ridge	Lasso
<b>MSSubClass</b>	-1.520608e-02	-0.012793	-0.016126
<b>LotFrontage</b>	7.817679e-03	0.002785	0.002086
<b>LotArea</b>	3.473573e-02	0.033040	0.024108
<b>OverallQual</b>	5.582310e-02	0.059251	0.066269
<b>OverallCond</b>	4.028924e-02	0.038615	0.040244
<b>YearRemodAdd</b>	2.369181e-02	0.023304	0.023800
<b>MasVnrArea</b>	1.167741e-05	-0.000431	0.000046
<b>BsmtFinSF1</b>	1.746055e+09	0.026051	0.015393
<b>BsmtFinSF2</b>	6.416146e+08	0.008115	0.000856
<b>BsmtUnfSF</b>	1.684956e+09	0.007080	-0.000000
<b>TotalBsmtSF</b>	-1.724279e+09	0.036318	0.037941
<b>1stFlrSF</b>	-1.062360e+08	0.047875	0.002400
<b>2ndFlrSF</b>	-1.161896e+08	0.048862	0.000000
<b>LowQualFinSF</b>	-1.286759e+07	-0.000116	-0.004057
<b>GrLivArea</b>	1.420053e+08	0.075785	0.142040
<b>BsmtFullBath</b>	1.230349e-02	0.015339	0.020087
<b>BsmtHalfBath</b>	6.692782e-04	0.000884	0.000000
<b>FullBath</b>	3.196806e-03	0.007554	0.006134
<b>HalfBath</b>	2.256468e-03	0.004803	0.004092
<b>BedroomAbvGr</b>	7.331699e-03	0.010011	0.004624
<b>KitchenAbvGr</b>	-1.431003e-02	-0.014394	-0.011783
<b>Fireplaces</b>	-5.295157e-03	-0.007101	0.000000
<b>GarageYrBlt</b>	4.395843e-03	0.000141	0.000000
<b>GarageCars</b>	2.910347e-02	0.033380	0.034627
<b>WoodDeckSF</b>	1.019961e-02	0.010473	0.009622
<b>OpenPorchSF</b>	4.070029e-03	0.003649	0.003829
<b>EnclosedPorch</b>	6.858930e-03	0.007358	0.003759
<b>3SsnPorch</b>	3.810376e-03	0.004740	0.003809
<b>ScreenPorch</b>	8.095803e-03	0.008703	0.007123
<b>PoolArea</b>	2.904532e-03	-0.002009	-0.004795
<b>MiscVal</b>	1.021191e-03	0.000939	-0.000000
<b>MoSold</b>	7.199794e-04	0.000438	-0.000000
<b>AgeOfHouse</b>	-6.239484e-02	-0.051905	-0.052985

	Linear	Ridge	Lasso
<b>MSZoning_FV</b>	9.274003e-02	0.068683	0.035276
<b>MSZoning_RH</b>	4.636907e-02	0.035282	0.018803
<b>MSZoning_RL</b>	1.793910e-01	0.134376	0.074449
<b>MSZoning_RM</b>	1.350354e-01	0.101120	0.046567
<b>Street_Pave</b>	5.741294e-03	0.006319	0.004758
<b>LotShape_IR2</b>	2.451539e-03	0.002675	0.000919
<b>LotShape_IR3</b>	-1.298539e-03	-0.003906	-0.004155
<b>LotShape_Reg</b>	2.470240e-03	0.002329	0.000000
<b>LandContour_HLS</b>	1.686316e-03	0.003427	0.000000
<b>LandContour_Low</b>	-6.874315e-04	0.000950	0.000000
<b>LandContour_Lvl</b>	7.216014e-03	0.009561	0.003703
<b>Utilities_NoSeWa</b>	-2.048017e-03	-0.002482	-0.002151
<b>LotConfig_CulDSac</b>	7.161677e-03	0.007142	0.007526
<b>LotConfig_FR2</b>	-7.268627e-03	-0.008260	-0.004452
<b>LotConfig_FR3</b>	-3.527999e-04	-0.001442	-0.000031
<b>LotConfig_Inside</b>	-2.848372e-03	-0.003431	-0.000060
<b>LandSlope_Mod</b>	4.582770e-03	0.004256	0.001350
<b>LandSlope_Sev</b>	-2.061029e-02	-0.017808	-0.008201
<b>Neighborhood_Blueste</b>	2.843235e-03	0.001326	0.000000
<b>Neighborhood_BrDale</b>	-7.559955e-04	-0.003419	-0.003551
<b>Neighborhood_BrkSide</b>	1.996967e-02	0.012328	0.006245
<b>Neighborhood_ClearCr</b>	1.058393e-02	0.009704	0.009280
<b>Neighborhood_CollgCr</b>	7.717423e-03	0.004379	0.000000
<b>Neighborhood_Crawfor</b>	3.352958e-02	0.030104	0.025709
<b>Neighborhood_Edwards</b>	-3.741689e-03	-0.010068	-0.012475
<b>Neighborhood_Gilbert</b>	3.471196e-03	0.000812	-0.001411
<b>Neighborhood_IDOTRR</b>	1.116958e-02	0.001513	-0.004760
<b>Neighborhood_MeadowV</b>	-1.194632e-02	-0.015819	-0.014356
<b>Neighborhood_Mitchel</b>	-4.804954e-04	-0.003327	-0.003740
<b>Neighborhood_NAmes</b>	1.242287e-02	0.006346	0.000000
<b>Neighborhood_NPkVill</b>	1.244128e-03	-0.000471	-0.000000
<b>Neighborhood_NWAmes</b>	1.604646e-03	-0.001293	-0.001193
<b>Neighborhood_NoRidge</b>	7.317960e-03	0.008112	0.005293

	Linear	Ridge	Lasso
<b>Neighborhood_NridgHt</b>	1.677114e-02	0.018370	0.016540
<b>Neighborhood_OldTown</b>	1.905072e-02	0.007900	-0.000000
<b>Neighborhood_SWISU</b>	1.289563e-02	0.009214	0.005156
<b>Neighborhood_Sawyer</b>	3.168374e-03	-0.000684	-0.001693
<b>Neighborhood_SawyerW</b>	6.479919e-03	0.004328	0.000061
<b>Neighborhood_Somerst</b>	1.994295e-02	0.019751	0.018077
<b>Neighborhood_StoneBr</b>	1.726687e-02	0.016214	0.012149
<b>Neighborhood_Timber</b>	2.365030e-03	0.001086	0.000000
<b>Neighborhood_Veenker</b>	8.448057e-03	0.007570	0.005161
<b>Condition1_Feedr</b>	1.464710e-02	0.011893	0.001218
<b>Condition1_Norm</b>	3.499979e-02	0.033646	0.020396
<b>Condition1_PosA</b>	7.450961e-03	0.007431	0.003023
<b>Condition1_PosN</b>	8.133732e-03	0.007695	0.002889
<b>Condition1_RRAe</b>	-3.750399e-04	-0.000768	-0.003997

```
In [78]: betas[betas['Lasso'] == 0].shape
```

```
Out[78]: (73, 3)
```

73 features have been removed by lasso

```
In [79]: # selected features :
betas.loc[betas['Lasso'] != 0, 'Lasso']
```

```

Out[79]: MSubClass          -0.016126
         LotFrontage       0.002086
         LotArea           0.024108
         OverallQual       0.066269
         OverallCond       0.040244
         YearRemodAdd      0.023800
         MasVnrArea        0.000046
         BsmtFinSF1        0.015393
         BsmtFinSF2        0.000856
         TotalBsmtSF       0.037941
         1stFlrSF          0.002400
         LowQualFinSF      -0.004057
         GrLivArea         0.142040
         BsmtFullBath      0.020087
         FullBath          0.006134
         HalfBath          0.004092
         BedroomAbvGr      0.004624
         KitchenAbvGr     -0.011783
         GarageCars        0.034627
         WoodDeckSF        0.009622
         OpenPorchSF       0.003829
         EnclosedPorch     0.003759
         3SsnPorch         0.003809
         ScreenPorch       0.007123
         PoolArea          -0.004795
         AgeOfHouse        -0.052985
         MSZoning_FV       0.035276
         MSZoning_RH       0.018803
         MSZoning_RL       0.074449
         MSZoning_RM       0.046567
         Street_Pave       0.004758
         LotShape_IR2      0.000919
         LotShape_IR3     -0.004155
         LandContour_Lvl   0.003703
         Utilities_NoSewa  -0.002151
         LotConfig_CulDSac 0.007526
         LotConfig_FR2     -0.004452
         LotConfig_FR3     -0.000031
         LotConfig_Inside  -0.000060
         LandSlope_Mod     0.001350
         LandSlope_Sev     -0.008201
         Neighborhood_BrDale -0.003551
         Neighborhood_BrkSide 0.006245
         Neighborhood_ClearCr 0.009280
         Neighborhood_Crawfor 0.025709
         Neighborhood_Edwards -0.012475
         Neighborhood_Gilbert -0.001411
         Neighborhood_IDOTRR -0.004760
         Neighborhood_MeadowV -0.014356
         Neighborhood_Mitchel -0.003740
         Neighborhood_NWAmes -0.001193
         Neighborhood_NoRidge 0.005293
         Neighborhood_NridgHt 0.016540
         Neighborhood_SWISU  0.005156
         Neighborhood_Sawyer -0.001693
         Neighborhood_SawyerW 0.000061
         Neighborhood_Somerst 0.018077
         Neighborhood_StoneBr 0.012149
         Neighborhood_Veenker 0.005161
         Condition1_Feetr   0.001218

```

Condition1_Norm	0.020396
Condition1_PosA	0.003023
Condition1_PosN	0.002889
Condition1_RRAe	-0.003997
Condition1_RRAn	0.003502
Condition1_RRNn	0.002185
Condition2_Norm	0.000992
Condition2_PosA	0.005142
Condition2_PosN	-0.057745
Condition2_RRAe	-0.006113
Condition2_RRAn	-0.001451
Condition2_RRNn	0.001917
BldgType_Twnhs	-0.008913
BldgType_TwnhsE	-0.000236
HouseStyle_1.5Unf	0.000985
HouseStyle_1Story	-0.004063
HouseStyle_2.5Fin	-0.003817
RoofStyle_Gable	-0.001728
RoofStyle_Mansard	0.004576
RoofStyle_Shed	0.007733
RoofMatl_CompShg	0.242023
RoofMatl_Membran	0.052098
RoofMatl_Metal	0.050648
RoofMatl_Roll	0.048608
RoofMatl_Tar&Grv	0.162588
RoofMatl_WdShake	0.101795
RoofMatl_WdShngl	0.121258
Exterior1st_BrkComm	-0.007722
Exterior1st_BrkFace	0.012118
Exterior1st_CBlock	-0.000181
Exterior1st_HdBoard	-0.000181
Exterior1st_ImStucc	0.000355
Exterior1st_MetalSd	0.003907
Exterior1st_Stone	0.000254
Exterior1st_VinylSd	0.010683
Exterior1st_Wd Sdng	-0.003691
Exterior1st_WdShing	0.002873
Exterior2nd_CBlock	-0.000016
Exterior2nd_CmentBd	0.006277
Exterior2nd_MetalSd	0.001039
Exterior2nd_Stucco	-0.001049
Exterior2nd_Wd Shng	-0.005511
MasVnrType_None	-0.001854
MasVnrType_Stone	0.000288
ExterQual_Fa	-0.001799
ExterQual_Gd	0.001780
ExterQual_TA	-0.000381
ExterCond_Fa	-0.002246
ExterCond_TA	0.004378
Foundation_PConc	0.008050
Foundation_Stone	0.002881
Foundation_Wood	-0.001469
BsmtQual_Gd	-0.016861
BsmtQual_TA	-0.015328
BsmtCond_Gd	0.002825
BsmtCond_Po	0.001494
BsmtCond_TA	0.005428
BsmtExposure_Gd	0.014815
BsmtExposure_No	-0.004859
BsmtExposure_NoBasement	-0.000684



BsmtFinType1_GLQ	0.007936
BsmtFinType1_LwQ	-0.001041
BsmtFinType1_Rec	-0.004179
BsmtFinType1_Unf	-0.008479
BsmtFinType2_BLQ	-0.007059
BsmtFinType2_GLQ	0.002634
BsmtFinType2_NoBasement	-0.003595
Heating_GasW	0.004351
Heating_Grav	-0.006446
Heating_OthW	-0.002529
Heating_Wall	0.001328
HeatingQC_Gd	-0.003212
HeatingQC_Po	-0.000592
HeatingQC_TA	-0.009003
CentralAir_Y	0.015706
Electrical_FuseF	0.003246
Electrical_FuseP	-0.001073
KitchenQual_Fa	-0.003111
KitchenQual_Gd	-0.019159
KitchenQual_TA	-0.015147
Functional_Maj2	-0.009753
Functional_Min2	0.002116
Functional_Mod	-0.004078
Functional_Sev	-0.009443
Functional_Typ	0.009887
FireplaceQu_Gd	0.003622
FireplaceQu_NoFireplace	-0.015091
FireplaceQu_Po	-0.002107
GarageType_Attchd	0.004673
GarageType_Basment	-0.001224
GarageType_Detchd	-0.000342
GarageFinish_Unf	-0.002298
GarageQual_Fa	-0.007903
GarageQual_Gd	0.003368
GarageCond_Fa	-0.003308
GarageCond_Po	0.001915
GarageCond_TA	0.000925
PavedDrive_P	0.004332
PavedDrive_Y	0.004804
SaleType_CWD	0.002690
SaleType_Con	0.002288
SaleType_ConLD	0.015459
SaleType_ConLw	0.000151
SaleType_New	0.024485
SaleType_Oth	0.000764
SaleCondition_AdjLand	0.003552
SaleCondition_Alloca	0.001206
SaleCondition_Family	0.002457
SaleCondition_Normal	0.017975
SaleCondition_Partial	0.000052

Name: Lasso, dtype: float64

## Top 10 features in Ridge Regression

```
In [80]: betas['Ridge'].sort_values(ascending=False)[:10]
```

```
Out[80]: RoofMatl_CompShg    0.275774
RoofMatl_Tar&Grv    0.183106
RoofMatl_WdShngl    0.137018
MSZoning_RL    0.134376
RoofMatl_WdShake    0.113699
MSZoning_RM    0.101120
GrLivArea    0.075785
MSZoning_FV    0.068683
RoofMatl_Membran    0.062715
RoofMatl_Metal    0.059642
Name: Ridge, dtype: float64
```

## Top 10 features in Lasso Regression

```
In [81]: betas['Lasso'].sort_values(ascending=False)[:10]
```

```
Out[81]: RoofMatl_CompShg    0.242023
RoofMatl_Tar&Grv    0.162588
GrLivArea    0.142040
RoofMatl_WdShngl    0.121258
RoofMatl_WdShake    0.101795
MSZoning_RL    0.074449
OverallQual    0.066269
RoofMatl_Membran    0.052098
RoofMatl_Metal    0.050648
RoofMatl_Roll    0.048608
Name: Lasso, dtype: float64
```

Optimal value of alpha in Ridge : 4

Optimal value of alpha in Lasso : 0.001

## Coding Questions:

### 1. Double the alpha

```
In [82]: #Fitting Ridge model for alpha = 4*2 i.e. 8 and printing coefficients
alpha = 8
ridge = Ridge(alpha=alpha)

ridge.fit(X_train, y_train)
print(ridge.coef_)
```

[-1.22268357e-02 2.75892347e-04 3.17078121e-02 6.09779876e-02  
3.76089785e-02 2.32102549e-02 -6.00220501e-04 2.02992812e-02  
7.38280273e-03 6.46270672e-03 2.96181397e-02 4.71718833e-02  
4.73560498e-02 5.74302843e-04 7.40889216e-02 1.70575483e-02  
9.28368717e-04 1.01561609e-02 6.27359029e-03 1.14705211e-02  
-1.42448830e-02 -7.38936885e-03 -1.94686713e-03 3.55508558e-02  
1.06230273e-02 3.34489675e-03 7.50200865e-03 5.15476054e-03  
8.99706167e-03 -4.72046655e-03 8.69967735e-04 2.62549776e-04  
-4.58324933e-02 5.46098945e-02 2.87962705e-02 1.08055384e-01  
8.08319712e-02 6.54462437e-03 2.74975166e-03 -5.48671075e-03  
2.17502100e-03 4.35745253e-03 1.84176784e-03 1.07569590e-02  
-2.70578682e-03 7.20684161e-03 -8.75369298e-03 -2.06419942e-03  
-3.71482461e-03 4.06216312e-03 -1.59765756e-02 5.47157372e-04  
-4.65239327e-03 8.57637531e-03 9.51160856e-03 2.90901461e-03  
2.85130651e-02 -1.31338441e-02 -4.32405024e-04 -3.46975569e-03  
-1.76720117e-02 -4.59208371e-03 3.55864167e-03 -1.24213417e-03  
-2.41688148e-03 8.87150527e-03 1.95830877e-02 2.23792618e-03  
7.51456616e-03 -2.43641600e-03 3.40569816e-03 1.99680594e-02  
1.58328564e-02 6.71212342e-04 7.23436472e-03 9.90426690e-03  
3.22838332e-02 7.25347556e-03 7.33596173e-03 -1.12086966e-03  
9.09207071e-03 1.44113995e-03 5.11751478e-03 1.40829199e-03  
1.96275371e-03 7.17069469e-03 -5.59116128e-02 -1.39218689e-02  
-1.94299276e-03 3.87890926e-03 -2.34990785e-03 -7.44363484e-04  
-9.04897202e-03 -3.06154749e-03 4.86175751e-04 -1.23045219e-02  
-4.66370512e-03 -4.60620656e-04 -1.06223935e-02 -2.30608397e-03  
-3.86088408e-03 -1.94280137e-02 -3.39863221e-03 -1.58235278e-02  
3.79188308e-03 1.54673610e-02 2.17453292e-01 4.99417873e-02  
4.75741119e-02 4.51055077e-02 1.44053106e-01 8.74171827e-02  
1.09204831e-01 2.15242003e-04 -7.48431770e-03 1.98460204e-02  
1.37695250e-04 -3.75473542e-03 1.01772033e-02 1.91376101e-03  
1.30622748e-02 9.33106958e-03 4.65420662e-03 3.40314473e-03  
2.57074052e-02 -1.07152971e-03 7.65243904e-03 1.24355291e-04  
2.47368499e-03 -2.70342060e-03 1.37695250e-04 2.01664723e-02  
3.53937911e-03 1.02249102e-03 8.33940152e-03 0.00000000e+00  
3.39768023e-03 -4.09886431e-03 1.24374708e-04 7.85831594e-03  
1.04108279e-02 -4.26910391e-03 1.26858036e-03 -2.33245071e-03  
1.38894247e-03 -2.85351856e-05 1.13063650e-02 9.16966829e-03  
-5.79858503e-03 -7.12201253e-03 0.00000000e+00 -2.16201446e-03  
1.22080474e-02 1.63405432e-02 -4.12279773e-03 3.20809781e-03  
-1.44800838e-03 -2.69734855e-03 -2.72772289e-02 3.35716867e-03  
-2.70870159e-02 5.93017368e-03 3.35716867e-03 8.92721971e-03  
8.12261028e-03 1.23818860e-02 -1.69056236e-03 -6.71731878e-03  
-7.99086981e-03 -3.69432845e-03 6.44714524e-03 -5.99299427e-03  
3.35716867e-03 -8.10174005e-03 -1.46847504e-02 -1.20187865e-02  
-2.68972937e-04 -4.96335213e-03 -1.18795981e-02 -6.26884586e-03  
-5.80431916e-03 1.55612180e-02 1.74308587e-02 -2.80085074e-03  
5.40853493e-04 9.06895649e-03 6.46951796e-04 -5.56440866e-03  
-3.45943425e-03 -1.20246546e-02 1.41288859e-02 5.53042990e-03  
2.57010677e-05 0.00000000e+00 -1.06212860e-03 -8.07780134e-03  
-3.26582852e-02 -2.89757146e-02 -9.36677784e-03 5.01111100e-03  
8.34009769e-03 -3.71541441e-03 -9.84773670e-03 1.89489977e-02  
6.83712518e-04 9.45510643e-03 -1.37968270e-02 -1.88813031e-03  
8.07657857e-03 4.16378539e-03 -1.64151795e-03 -3.58525400e-05  
3.75358304e-03 -2.23541856e-03 -3.25878248e-03 -3.25878248e-03  
-1.38805945e-03 -4.89592282e-03 -1.49599279e-02 3.51011748e-03  
-3.25878248e-03 -2.62979174e-03 -5.17363567e-03 -8.78926870e-03  
-3.52453972e-03 -3.25878248e-03 4.47104295e-03 -9.34922825e-03  
5.68878070e-03 6.17906912e-03 3.48249055e-03 4.30756583e-03  
1.90268754e-02 4.54455503e-04 1.22264470e-03 1.41150158e-02

```
2.68348661e-03  3.26133828e-03  4.27890349e-03  5.80804059e-03
4.67909410e-03  1.99649841e-02  1.39602663e-02]
```

In [83]: *# Lets calculate some metrics such as R2 score, RSS and RMSE*

```
y_pred_train = ridge.predict(X_train)
y_pred_test = ridge.predict(X_test)

metric2 = []
r2_train_lr = r2_score(y_train, y_pred_train)
print("R square train :",r2_train_lr)
metric2.append(r2_train_lr)

r2_test_lr = r2_score(y_test, y_pred_test)
print("R square test :",r2_test_lr)
metric2.append(r2_test_lr)

rss1_lr = np.sum(np.square(y_train - y_pred_train))
print("Rss train :",rss1_lr)
metric2.append(rss1_lr)

rss2_lr = np.sum(np.square(y_test - y_pred_test))
print("Rss test :",rss2_lr)
metric2.append(rss2_lr)

mse_train_lr = mean_squared_error(y_train, y_pred_train)
print("MSE train:",mse_train_lr)
metric2.append(mse_train_lr)

print("RMSE train:",mse_train_lr**0.5)
metric2.append(mse_train_lr**0.5)

mse_test_lr = mean_squared_error(y_test, y_pred_test)
print("MSE test :",mse_test_lr)
metric2.append(mse_test_lr)

print("RMSE test :",mse_test_lr**0.5)
metric2.append(mse_test_lr**0.5)
```

```
R square train : 0.9537251879652039
R square test : 0.8721002342302973
Rss train : 7.426653074278145
Rss test : 9.217430166567063
MSE train: 0.007273901150125509
RMSE train: 0.08528716873085605
MSE test : 0.021044361110883706
RMSE test : 0.1450667470886547
```

In [84]: `lasso = Lasso(alpha=0.002)`

```
lasso.fit(X_train, y_train)
print(lasso.coef_)
```

[-1.81729275e-02 0.00000000e+00 1.79277302e-02 7.74425269e-02  
3.87982651e-02 2.25123786e-02 0.00000000e+00 4.67153891e-03  
0.00000000e+00 -0.00000000e+00 2.69621334e-02 2.57148749e-03  
0.00000000e+00 -2.47926646e-03 1.37906627e-01 2.51074610e-02  
0.00000000e+00 7.70660788e-03 3.60084624e-03 4.50337940e-03  
-9.19304934e-03 0.00000000e+00 -0.00000000e+00 3.89809306e-02  
8.96031751e-03 1.63425190e-03 2.42212987e-03 2.68895257e-03  
6.68041484e-03 -9.57214431e-03 -0.00000000e+00 -0.00000000e+00  
-4.78551138e-02 5.01167552e-03 4.60324406e-03 2.06749814e-02  
0.00000000e+00 4.03643376e-03 4.70546249e-04 -6.91927177e-03  
0.00000000e+00 0.00000000e+00 0.00000000e+00 2.63176396e-03  
-2.22199809e-03 7.76956318e-03 -2.96609919e-03 -0.00000000e+00  
-0.00000000e+00 0.00000000e+00 -0.00000000e+00 -0.00000000e+00  
-2.69510428e-03 3.35017403e-03 1.08748521e-02 0.00000000e+00  
2.50871647e-02 -1.22856070e-02 -3.52464177e-04 -9.53433420e-03  
-1.15827327e-02 -1.72562306e-03 0.00000000e+00 -0.00000000e+00  
-0.00000000e+00 6.89846544e-03 2.03410448e-02 -1.83794774e-03  
3.05061755e-03 -1.00883910e-03 0.00000000e+00 2.06573452e-02  
1.03781719e-02 0.00000000e+00 4.61417406e-03 -0.00000000e+00  
1.87315016e-02 1.81775880e-03 2.00872944e-03 -3.37937654e-03  
2.87473979e-03 0.00000000e+00 1.40907286e-03 -0.00000000e+00  
2.63014080e-04 3.69358812e-03 -5.15556893e-02 -1.44177848e-06  
-5.32293402e-04 0.00000000e+00 0.00000000e+00 -0.00000000e+00  
-1.04596812e-02 -6.13790994e-04 3.32144207e-05 -0.00000000e+00  
-1.61177628e-03 -0.00000000e+00 -0.00000000e+00 -0.00000000e+00  
-0.00000000e+00 -2.27953513e-03 0.00000000e+00 0.00000000e+00  
3.42816278e-03 2.22002840e-04 1.27549406e-01 2.48032041e-02  
2.46098825e-02 2.45774608e-02 8.32408493e-02 5.22931979e-02  
6.65765980e-02 -0.00000000e+00 -7.88907980e-03 1.09862051e-02  
-0.00000000e+00 0.00000000e+00 -1.57675981e-03 0.00000000e+00  
1.12666659e-03 -0.00000000e+00 0.00000000e+00 -1.44017786e-03  
5.75479865e-03 -5.30971117e-03 0.00000000e+00 -0.00000000e+00  
-0.00000000e+00 0.00000000e+00 -0.00000000e+00 3.86881622e-03  
-0.00000000e+00 0.00000000e+00 0.00000000e+00 0.00000000e+00  
-0.00000000e+00 0.00000000e+00 -3.69354277e-03 0.00000000e+00  
-0.00000000e+00 -5.13181190e-03 0.00000000e+00 -0.00000000e+00  
0.00000000e+00 -5.01913992e-03 0.00000000e+00 -3.83620388e-03  
-1.62091502e-03 -0.00000000e+00 0.00000000e+00 2.34822021e-03  
0.00000000e+00 8.88572800e-03 -0.00000000e+00 1.70728706e-03  
-0.00000000e+00 0.00000000e+00 -1.15615385e-02 -0.00000000e+00  
-8.73746264e-03 1.45194190e-03 -0.00000000e+00 0.00000000e+00  
5.10344149e-03 1.66491439e-02 -0.00000000e+00 -4.70605819e-03  
-2.59199501e-03 -0.00000000e+00 8.69301957e-03 -0.00000000e+00  
-0.00000000e+00 -3.37195017e-03 -1.22854754e-02 -6.18678611e-03  
1.68266989e-03 -0.00000000e+00 -2.72305489e-03 -0.00000000e+00  
0.00000000e+00 -0.00000000e+00 5.50156850e-03 -3.83825505e-03  
-2.43876551e-03 1.49085518e-04 0.00000000e+00 -1.34770746e-03  
-0.00000000e+00 -8.31609179e-03 1.63422699e-02 2.02219763e-03  
-0.00000000e+00 0.00000000e+00 0.00000000e+00 -1.87276204e-04  
-8.00903310e-03 -5.20464973e-03 -8.94829883e-03 0.00000000e+00  
0.00000000e+00 -8.17352786e-04 -7.85293092e-03 7.00626913e-03  
0.00000000e+00 2.63590329e-03 -1.82933178e-02 -1.67084493e-03  
0.00000000e+00 4.32253678e-03 -4.44889518e-04 0.00000000e+00  
-0.00000000e+00 -0.00000000e+00 -0.00000000e+00 -0.00000000e+00  
-0.00000000e+00 -4.03065630e-03 -6.79922091e-03 2.69536145e-03  
-0.00000000e+00 0.00000000e+00 0.00000000e+00 -1.64119341e-03  
-0.00000000e+00 -0.00000000e+00 0.00000000e+00 0.00000000e+00  
3.63965538e-03 3.37453105e-03 2.10136440e-03 1.61987324e-03  
1.10780478e-02 -0.00000000e+00 0.00000000e+00 2.25643001e-02

```
0.00000000e+00 0.00000000e+00 9.70378752e-04 -0.00000000e+00
8.52051642e-04 1.55748089e-02 0.00000000e+00]
```

In [85]: *# Lets calculate some metrics such as R2 score, RSS and RMSE*

```
y_pred_train = lasso.predict(X_train)
y_pred_test = lasso.predict(X_test)

metric3 = []
r2_train_lr = r2_score(y_train, y_pred_train)
print("R square train :",r2_train_lr)
metric3.append(r2_train_lr)

r2_test_lr = r2_score(y_test, y_pred_test)
print("R square test :",r2_test_lr)
metric3.append(r2_test_lr)

rss1_lr = np.sum(np.square(y_train - y_pred_train))
print("Rss train :",rss1_lr)
metric3.append(rss1_lr)

rss2_lr = np.sum(np.square(y_test - y_pred_test))
print("Rss test :",rss2_lr)
metric3.append(rss2_lr)

mse_train_lr = mean_squared_error(y_train, y_pred_train)
print("MSE train:",mse_train_lr)
metric3.append(mse_train_lr)

print("RMSE train:",mse_train_lr**0.5)
metric3.append(mse_train_lr**0.5)

mse_test_lr = mean_squared_error(y_test, y_pred_test)
print("MSE test :",mse_test_lr)
metric3.append(mse_test_lr)

print("RMSE test :",mse_test_lr**0.5)
metric3.append(mse_test_lr**0.5)
```

```
R square train : 0.938111442795243
R square test : 0.8714626375913678
Rss train : 9.932505901520138
Rss test : 9.26338023112271
MSE train: 0.009728213419706306
RMSE train: 0.09863170595557144
MSE test : 0.021149269934070115
RMSE test : 0.14542788568245815
```

In [86]: *# Creating a table which contain all the metrics*

```
lr_table = {'Metric': ['R2 Score (Train)','R2 Score (Test)','RSS (Train)','RSS (Test)']

lr_metric = pd.DataFrame(lr_table ,columns = ['Metric'] )

rg_metric = pd.Series(metric2, name = 'Ridge Regression')
ls_metric = pd.Series(metric3, name = 'Lasso Regression')

final_metric = pd.concat([lr_metric, rg_metric, ls_metric], axis = 1)


final_metric
```

Out[86]:

	Metric	Ridge Regression	Lasso Regression
0	R2 Score (Train)	0.953725	0.938111
1	R2 Score (Test)	0.872100	0.871463
2	RSS (Train)	7.426653	9.932506
3	RSS (Test)	9.217430	9.263380
4	MSE (Train)	0.007274	0.009728
5	MSE (Test)	0.085287	0.098632
6	RMSE (Train)	0.021044	0.021149
7	RMSE (Test)	0.145067	0.145428

In [87]:

```
#plotting
model_cv_results = pd.DataFrame(model_cv.cv_results_)
plt.plot(model_cv_results['param_alpha'], model_cv_results['mean_train_score'])
plt.plot(model_cv_results['param_alpha'], model_cv_results['mean_test_score'])
plt.xlabel('alpha')
plt.xscale('log')
plt.ylabel('R2 value')
plt.legend(['train score', 'test score'])
plt.show()
```



In [88]:

```
betas = pd.DataFrame(index=X.columns)
betas.rows = X.columns
betas['Linear'] = lm.coef_
betas['Ridge'] = ridge.coef_
betas['Lasso'] = lasso.coef_
```

```
In [89]: betas['Ridge'].sort_values(ascending=False)[:10]
```

```
Out[89]: RoofMatl_CompShg      0.217453
RoofMatl_Tar&Grv      0.144053
RoofMatl_WdShngl      0.109205
MSZoning_RL           0.108055
RoofMatl_WdShake      0.087417
MSZoning_RM           0.080832
GrLivArea             0.074089
OverallQual           0.060978
MSZoning_FV           0.054610
RoofMatl_Membran      0.049942
Name: Ridge, dtype: float64
```

```
In [90]: betas['Lasso'].sort_values(ascending=False)[:10]
```

```
Out[90]: GrLivArea          0.137907
RoofMatl_CompShg      0.127549
RoofMatl_Tar&Grv      0.083241
OverallQual           0.077443
RoofMatl_WdShngl      0.066577
RoofMatl_WdShake      0.052293
GarageCars            0.038981
OverallCond           0.038798
TotalBsmtSF           0.026962
BsmtFullBath          0.025107
Name: Lasso, dtype: float64
```

### 3. Removing the top 5 variables

```
In [91]: # top 5 variables are : RoofMatl_CompShg, RoofMatl_Tar&Grv, GrLivArea, RoofMatl_WdShngl
# storing them in a variable
top5 = ['RoofMatl_CompShg', 'RoofMatl_Tar&Grv', 'GrLivArea', 'RoofMatl_WdShngl', 'RoofMatl_WdShake']
```

```
Out[91]: ['RoofMatl_CompShg',
'RoofMatl_Tar&Grv',
'GrLivArea',
'RoofMatl_WdShngl',
'RoofMatl_WdShake']
```

```
In [92]: # printing X train and test set
X_train.head()
```



```
Out[92]:
```

	MSSubClass	LotFrontage	LotArea	OverallQual	OverallCond	YearRemodAdd	MasVnrArea	B
<b>210</b>	-0.657071	-0.115302	-0.473765	-0.779861	0.383154	-1.694350	-0.558025	
<b>318</b>	0.035976	0.926898	-0.056845	0.649651	-0.533005	0.390956	0.809137	
<b>239</b>	-0.195040	-0.794998	-0.169324	-0.065105	-1.449164	-1.694350	-0.558025	
<b>986</b>	-0.195040	-0.477806	-0.502297	-0.065105	2.215472	0.875911	-0.558025	
<b>1416</b>	3.039179	-0.432493	0.082905	-1.494617	0.383154	-1.694350	-0.558025	

5 rows × 243 columns

```
In [93]: X_test.head()
```

```
Out[93]:
```

	MSSubClass	LotFrontage	LotArea	OverallQual	OverallCond	YearRemodAdd	MasVnrArea	B
<b>1436</b>	-0.888086	-0.432493	-0.144189	-1.494617	0.383154	-0.675945	-0.558025	
<b>57</b>	0.035976	0.881585	0.112505	0.649651	-0.533005	0.924407	-0.558025	
<b>780</b>	-0.888086	-0.296554	-0.253368	0.649651	-0.533005	0.536443	-0.355087	
<b>382</b>	0.035976	0.428455	-0.120412	0.649651	-0.533005	1.021398	-0.558025	
<b>1170</b>	0.498007	0.292515	-0.058786	-0.065105	0.383154	-0.384972	-0.558025	

5 rows × 243 columns

```
In [94]: X_train1 = X_train.drop(top5, axis = 1)
X_test1 = X_test.drop(top5, axis = 1)
```

```
In [95]: lasso = Lasso()

# cross validation
model_cv = GridSearchCV(estimator = lasso,
                        param_grid = params,
                        scoring= 'neg_mean_absolute_error',
                        cv = folds,
                        return_train_score=True,
                        verbose = 1)

model_cv.fit(X_train1, y_train)
```

Fitting 5 folds for each of 28 candidates, totalling 140 fits

```
Out[95]:
```

```
GridSearchCV
  estimator: Lasso
    Lasso
```

```
In [96]: model_cv.best_params_
```

Out[96]: {'alpha': 0.001}

```
In [97]: #Fitting Lasso model for alpha = 0.001 and printing coefficients which have been penal

alpha =0.001

lasso = Lasso(alpha=alpha)

lasso.fit(X_train1, y_train)
```

Out[97]:

▼

Lasso

Lasso(alpha=0.001)

```
In [98]: # Lets calculate some metrics such as R2 score, RSS and RMSE
```

```
y_pred_train = lasso.predict(X_train1)
y_pred_test = lasso.predict(X_test1)

metric3 = []
r2_train_lr = r2_score(y_train, y_pred_train)
print("R square train :",r2_train_lr)
metric3.append(r2_train_lr)

r2_test_lr = r2_score(y_test, y_pred_test)
print("R square test :",r2_test_lr)
metric3.append(r2_test_lr)

rss1_lr = np.sum(np.square(y_train - y_pred_train))
print("Rss train :",rss1_lr)
metric3.append(rss1_lr)

rss2_lr = np.sum(np.square(y_test - y_pred_test))
print("Rss test :",rss2_lr)
metric3.append(rss2_lr)

mse_train_lr = mean_squared_error(y_train, y_pred_train)
print("MSE train:",mse_train_lr)
metric3.append(mse_train_lr)

print("RMSE train:",mse_train_lr**0.5)
metric3.append(mse_train_lr**0.5)

mse_test_lr = mean_squared_error(y_test, y_pred_test)
print("MSE test :",mse_test_lr)
metric3.append(mse_test_lr)

print("RMSE test :",mse_test_lr**0.5)
metric3.append(mse_test_lr**0.5)
```

```
R square train : 0.9332528937199062
R square test : 0.8710398117759115
Rss train : 10.712255334100195
Rss test : 9.293852276189687
MSE train: 0.010491924910969829
RMSE train: 0.10243009768114951
MSE test : 0.021218840813218464
RMSE test : 0.1456668830352955
```

```
In [99]: betas = pd.DataFrame(index=X_train1.columns)
betas.rows = X_train1.columns

betas['Lasso'] = lasso.coef_
```

```
In [100... betas['Lasso'].sort_values(ascending=False)[:10]
```

```
Out[100]: 2ndFlrSF          0.100230
1stFlrSF          0.099158
MSZoning_RL       0.084011
OverallQual       0.072898
MSZoning_RM       0.057104
GarageCars        0.041877
MSZoning_FV       0.039181
OverallCond       0.033899
Neighborhood_Crawfor 0.027526
BsmtFullBath      0.025117
Name: Lasso, dtype: float64
```

```
In [ ]:
```